

ENGLISH FOR WORK / SEPTEMBER 2026

AI Development English

Conversation lab

10 complete workplace dialogues, followed by eight additional role-play cases. Study the exchanges, then make the conversation your own.

PRACTICE THAT TRANSFERS TO WORK

Read a realistic case. Find the right language. Speak, write, get feedback, and try again.

Eight lessons | Intermediate to advanced | Independent or classroom study

For AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders.

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

All scenarios, figures, and model responses are fictional teaching material. Use invented details in practice.

[Open this course on English Ladder](#)

Inside this conversation lab

Original fictional training conversations in professional English. These are realistic scripts, not transcripts of actual people. The professional roles are the speaker labels; all figures, incidents, and organizations are invented.

- 01 Standup: RAG Latency Spike
- 02 Design Review: Prompt Fix or Fine-Tune?
- 03 Data Meeting: Label Ambiguity
- 04 Incident Update: Tool-Calling Failure
- 05 Eval Readout: Better Average, Worse Edge Cases
- 06 Security Review: Prompt Injection
- 07 Customer Call: Hallucination Report
- 08 Product Planning: Model Choice
- 09 Research Sync: Benchmark vs Product Eval
- 10 Executive Briefing: Risk and Confidence

Then try eight independent role-play cases, followed by feedback criteria and references.

Read it. Hear it. Make it yours.

1. Read the setting. Identify the roles, the immediate problem, and what each speaker needs.
2. Read the whole conversation aloud with a partner. In a group, give a third person the observer role. For three-speaker scripts, assign each professional a separate voice.
3. Notice how the speakers ask a specific question, qualify a claim, disagree, explain a technical point, or agree on a next step. Underline language you would actually use.
4. Cover the script. Recreate the exchange using the situation and professional roles. Keep the meaning, but change the wording.
5. Switch roles and introduce a complication: a missing record, a different priority, an unavailable colleague, or a changed assumption. End with a clear outcome or unresolved question.

Register and terminology

These colleagues use concise professional language. In an internal technical meeting, familiar abbreviations can be natural. With a new colleague or a client, expand unfamiliar terms and check understanding. Keep disagreement directed at the evidence or proposal, and match the degree of certainty to what is known.

Independent study

Speak each role aloud. Pause before the reply and supply your own response before revealing the next line. Record a second version using invented details, then compare its clarity and tone with the script.

Professional context

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

FULL DIALOGUE 01 / 7 SPEAKING TURNS

Standup: RAG Latency Spike

Morning standup for a document assistant.

Roles: PM / Engineer / AI technical lead

PM: Are we still on track for the pilot on Friday?

Engineer: Functionally, yes. The blocker is latency. P95 went from 4.2 seconds to 8.7 after we added reranking.

AI technical lead: So the answer quality improved, but the serving path is too slow. Is the bottleneck retrieval, reranking, or generation?

Engineer: Mostly reranking. The generator time is stable.

AI technical lead: Then my proposal is to keep reranking for high-risk queries only and use the faster path for simple FAQ queries. I can bring an ablation by end of day.

PM: Include the quality tradeoff for each query class, not only the average latency.

AI technical lead: I will compare like-for-like traffic and document the routing assumptions before recommending a pilot configuration.

Notice the professional language

P95 means the 95th percentile request latency.

Ablation means testing the effect of removing or changing one component.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 02 / 7 SPEAKING TURNS

Design Review: Prompt Fix or Fine-Tune?

Design review for a customer-support summarizer.

Roles: Researcher / Product manager / AI technical lead

Researcher: We should fine-tune. The model keeps missing refund-policy exceptions.

Product manager: Would fine-tuning actually solve that, or is the policy changing too often?

AI technical lead: I would separate style from facts. If the issue is current policy knowledge, RAG may be safer. If the issue is summary format, a prompt or fine-tune could help.

Researcher: Good point. The examples show both problems.

AI technical lead: Let's run two eval slices: factual policy coverage and format compliance. Then we can choose the adaptation method with evidence.

Researcher: I will label the examples by failure type so we can compare the two eval slices.

AI technical lead: Keep the policy version fixed for that comparison; otherwise the adaptation results will be hard to interpret.

Notice the professional language

RAG - Retrieval-augmented generation: retrieve relevant context, then generate an answer using it.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 03 / 6 SPEAKING TURNS

Data Meeting: Label Ambiguity

Annotation guideline meeting for safety classification.

Roles: Data lead / AI technical lead / Reviewer

Data lead: Annotators disagree on whether this is medical advice or general wellness information.

AI technical lead: For this project, we need an agreed annotation rule and worked examples that distinguish the categories. Let's have the domain reviewers resolve the boundary cases before treating the labels as reliable ground truth.

Reviewer: What about borderline cases?

AI technical lead: We should add a borderline tag and route those to expert review. Otherwise the ground truth will be noisy.

Data lead: Who resolves cases where the reviewers still disagree?

AI technical lead: Let's assign an adjudication owner and record the rationale, then test the revised guidance on a held-out sample.

Notice the professional language

Noisy labels are labels that are inconsistent, wrong, or ambiguous.

Ground truth should be as consistent as possible, but it is often a human-created artifact.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 04 / 8 SPEAKING TURNS

Incident Update: Tool-Calling Failure

Incident channel after an agent booked duplicate appointments.

Roles: Support / AI technical lead / Engineering manager

Support: Customers are seeing duplicate calendar events.

AI technical lead: We found the immediate cause. The model retried the booking tool after a timeout, but the first call had actually succeeded.

Engineering manager: Mitigation?

AI technical lead: We paused the affected workflow and are reviewing retry handling and idempotency with the responsible engineers. We need to verify the complete fix, including how duplicate requests are identified, before resuming it.

Support: What should we tell customers?

AI technical lead: Say we are investigating duplicate events and will provide a verified update. Any correction or deletion needs the appropriate user authorization and checks; we shouldn't promise that records have already been repaired.

Support: I will identify affected records without changing them until the correction plan is authorized.

AI technical lead: Good. Preserve the incident evidence and distinguish confirmed duplicates from records that only look similar.

Notice the professional language

Non-idempotent means repeating the action can create a different or duplicate result.

An idempotency key helps the system recognize that a retry belongs to the same intended action.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 05 / 8 SPEAKING TURNS

Eval Readout: Better Average, Worse Edge Cases

Release meeting for a model upgrade.

Roles: PM / AI technical lead / Researcher

PM: The average score is up. Can we ship?

AI technical lead: I recommend a hold. The overall pass rate improved from 86% to 89%, but the legal-disclaimer slice regressed by 11 points.

Researcher: Is that statistically meaningful?

AI technical lead: The sample is small, so I would not overclaim. But the failures are severe enough to block release until we inspect them.

PM: What is the next step?

AI technical lead: Human review of the failed slice today, then a targeted prompt or policy fix before we rerun the regression suite.

PM: Keep the launch decision pending and show the severe failures at the review.

AI technical lead: I will include the sample size, criteria, and rerun results so the release owner can evaluate the actual evidence.

Notice the professional language

Pass rate - The percentage of eval cases meeting the success criterion.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 06 / 6 SPEAKING TURNS

Security Review: Prompt Injection

Security review for a browsing agent.

Roles: Security engineer / AI technical lead / Developer

Security engineer: The page contains text telling the agent to ignore the system instructions.

AI technical lead: That is prompt injection from untrusted content. The model should treat page text as data, not instructions.

Developer: Can we just add a stronger system prompt?

AI technical lead: A stronger prompt helps, but it is not enough. We need tool permissions, allowlists, confirmation for risky actions, and logging for suspicious instructions.

Security engineer: I will map the tools and the authority each one has.

AI technical lead: Then we can test the trust boundaries and permission controls instead of relying on a prompt change as complete protection.

Notice the professional language

Prompt injection - Untrusted input tries to manipulate model instructions or tool use.

System prompt - High-priority instructions that guide model behavior inside an application.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 07 / 6 SPEAKING TURNS

Customer Call: Hallucination Report

Customer reports that the assistant invented a policy.

Roles: Customer / AI technical lead

Customer: The answer cited a policy that does not exist.

AI technical lead: Thank you. We should call that an unsupported answer, not a confirmed policy source. Can you share the prompt, output, and timestamp?

Customer: Yes. Does this mean the model is unreliable?

AI technical lead: It means our grounding failed in this case. We will check whether retrieval missed the right document, whether the prompt allowed unsupported claims, or whether the citation validator failed.

Customer: I can provide the example through your approved support channel.

AI technical lead: Please include only the information needed for the investigation. We will explain what failed and what the evidence supports.

Notice the professional language

Use calm, precise failure language with customers.

Do not blame the model before investigating retrieval, prompt, validation, and UI layers.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 08 / 8 SPEAKING TURNS

Product Planning: Model Choice

Planning meeting for a high-volume summarization feature.

Roles: Finance / AI technical lead / PM

Finance: The larger model is too expensive for this volume.

AI technical lead: We can test a smaller model with a stricter prompt and a post-generation validator. The question is whether quality remains above the release threshold.

PM: What would you measure?

AI technical lead: Summary faithfulness, key-point coverage, latency, cost per thousand requests, and human escalation rate.

Finance: Can you bring options?

AI technical lead: Yes. I will compare three paths: larger model, smaller model with validation, and hybrid routing for complex cases.

Finance: Show costs under the same traffic mix and quality threshold.

AI technical lead: I will include failure rates, escalation workload, and latency distributions so the comparison reflects the whole serving path.

Notice the professional language

Hybrid routing sends different requests to different models or paths based on complexity or risk.

Cost discussions should include quality and operational risk, not only token price.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 09 / 6 SPEAKING TURNS

Research Sync: Benchmark vs Product Eval

Research team proposes a model because it performs well on a public benchmark.

Roles: Researcher / AI technical lead

Researcher: This checkpoint is strong on the benchmark.

AI technical lead: That is promising, but the benchmark may not represent our user traffic. We need a product eval before we switch.

Researcher: What gap do you expect?

AI technical lead: Our users ask mixed-language, document-grounded questions with messy formatting. The public benchmark may not test retrieval grounding or citation quality.

Researcher: I will add representative messy documents and mixed-language questions to the test set.

AI technical lead: Check for evaluation leakage and document the sampling approach before interpreting the checkpoint comparison.

Notice the professional language

Benchmark strength is useful evidence, not a release decision by itself.

Product evals should resemble the actual distribution of user tasks.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 10 / 8 SPEAKING TURNS

Executive Briefing: Risk and Confidence

Briefing a VP before an AI feature launch.

Roles: VP / AI technical lead

VP: Are we confident enough to launch?

AI technical lead: We are confident for the internal beta, not for general availability. The main remaining risks are unsupported answers in long-tail documents and latency during peak usage.

VP: What controls are in place?

AI technical lead: We have source citations, a refusal path for low-confidence retrieval, human review for escalations, and daily sampling of traces.

VP: What would make you stop the beta?

AI technical lead: A severe unsupported answer, repeated privacy exposure, or P95 latency above ten seconds for more than one hour.

VP: Make those stop criteria explicit in the beta plan and identify the release owner.

AI technical lead: I will confirm the approved thresholds and response process; the numbers in this scenario are team-specific criteria, not universal safety standards.

Notice the professional language

Executives need decision-grade confidence, not all technical detail.

Name launch scope: internal beta, limited pilot, general availability, or rollback.

Rehearse and extend

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

Eight more situations to make your own

The following cases are open role plays. They give you facts, contrasting roles, useful expressions, and a short model response. Use what you learned from the complete dialogues to create a longer exchange.

Preparation: two minutes. Conversation: three to five minutes. Feedback: one minute. Switch roles and repeat with the second-round challenge.

ROLE-PLAY CASE 01

Speaking the AI Development Stack

SHARED CASE

A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Role A | Responding professional

Explain a useful distinction in plain English. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Useful expressions

In this context, ... means

The difference is that

How would you explain that distinction to a colleague?

Cover this until after your first attempt

There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

Second round

Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 02

LLMs, Transformers, Tokens, and Context

SHARED CASE

A prompt plus retrieved documents exceeds a model's context limit. A colleague proposes simply increasing the output length.

Role A | Responding professional

Clarify missing information before responding. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Useful expressions

When you say ..., do you mean ...?

Could you clarify which ...?

So, my understanding is Is that right?

Cover this until after your first attempt

Do you mean a longer answer or more room for the input? Those are different limits. Let's check which messages and documents must fit in the context window.

Second round

The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 03

Data, Datasets, Labels, and Leakage

SHARED CASE

An evaluation scores 94%, but some evaluation examples also appear in the training data. A release meeting starts in an hour.

Role A | Responding professional

Separate observations, assumptions, and conclusions. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to tell a colleague what is known. Ask which part is confirmed and which part is an assumption. Challenge one statement that sounds more certain than the case supports.

Useful expressions

The available evidence shows

This may indicate ..., but

We have not yet established whether

Cover this until after your first attempt

The 94% result may be inflated because the datasets overlap. It does not yet show performance on unseen examples. I'll label this result and request a separate evaluation set.

Second round

Someone asks for a yes-or-no answer when the evidence supports only a qualified answer. Be concise while preserving the uncertainty.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 04

Retrieval, Embeddings, Vector Search, and RAG

SHARED CASE

A retrieval experiment with reranking finds more relevant passages but adds 300 milliseconds to response time. No user study has been run.

Role A | Responding professional

Compare options using the same criteria and state the tradeoff. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You favor one option but have not checked whether the scope is comparable. Ask about one difference, then explain the tradeoff you are willing to accept.

Useful expressions

Both options ..., but

Option A ..., whereas option B

The comparison depends on

Cover this until after your first attempt

Reranking improved passage relevance in this test, but it added 300 milliseconds. We still need to check the effect on answer quality and the user experience before choosing a configuration.

Second round

Your listener focuses only on the lowest price or fastest date. Explain one other relevant comparison criterion.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 05

Fine-Tuning, Alignment, and Adaptation

SHARED CASE

A summarizer uses the wrong format and sometimes cites an outdated refund policy. A colleague proposes fine-tuning as one solution to both problems.

Role A | Responding professional

Clarify missing information before responding. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Useful expressions

When you say ..., do you mean ...?

Could you clarify which ...?

So, my understanding is Is that right?

Cover this until after your first attempt

Could we separate the formatting problem from the outdated information? We can test a prompt change for the format and an updated retrieval source for policy facts before choosing an adaptation method.

Second round

The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 06

Evaluation, Benchmarks, and Regression

SHARED CASE

A new model passes 88 of 100 test cases; the previous version passes 90. The new version performs better on long documents but worse on short requests.

Role A | Responding professional

Distinguish completed work from pending work and explain its impact. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to brief someone who has only one minute. Ask what is complete, what is open, and which open item affects the next action.

Useful expressions

We have completed ...; ... remains open.

As of ..., the status is

The next update will cover

Cover this until after your first attempt

The overall pass rate fell from 90% to 88%. Long-document results improved, while short-request results declined. I'll show both groups so the release decision reflects the tradeoff.

Second round

The listener asks you to reduce the update to two sentences. Preserve the most important completed and pending items.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 07

Inference, Latency, Cost, and Deployment

SHARED CASE

A team can reduce average response time from four seconds to two by using a smaller model. Complex-answer quality has not yet been checked.

Role A | Responding professional

Negotiate scope or timing while making conditions explicit. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You want the preferred outcome but can accept one tradeoff. Ask for two options, then state which condition you can accept and which still needs approval.

Useful expressions

We can ..., provided that

Would you prefer ... or ...?

I can take that proposal for review; I cannot yet confirm

Cover this until after your first attempt

The smaller model is faster in this test. If it meets our quality criteria on complex requests, we could use it more widely. Until then, I'd keep the rollout limited.

Second round

The preferred option is unavailable. Offer an alternative and clearly state any approval still needed.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 08

Safety, Security, Privacy, and Governance

SHARED CASE

A test assistant repeats a fictional secret embedded in a retrieved document. The team has a reporting channel and needs a concise reproduction note.

Role A | Responding professional

Make a request with a clear action, purpose, and realistic timing. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You can help but need a clear request and its purpose. Ask what is required and whether the timing is fixed or negotiable.

Useful expressions

Could you provide ... so that ...?

Please confirm ... by [agreed time].

If that timing is not possible, please let me know

Cover this until after your first attempt

Please review this test through the security reporting channel. I've attached the synthetic input, observed output, and reproduction steps. We have not yet established whether other configurations are affected.

Second round

The other person cannot meet the requested timing. Clarify an alternative without inventing authority to change a formal deadline.

Debrief

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

Give feedback that helps

Score each criterion 0, 1, or 2. This is a practice rubric, not a professional qualification or CEFR test. Do not award or remove points for a particular accent.

Meaning and accuracy

Keeps the case facts accurate; clearly separates confirmed and unknown information.

Clarity and language

Uses understandable sentences and explains specialist terms when the listener needs it.

Interaction and tone

Responds to the other person's question, checks understanding, and uses an appropriate tone.

Task completion

Makes the requested action or unresolved question clear without inventing authority or facts.

Use the same scale for each criterion

0 - Not yet: The reader or listener cannot yet recover the intended meaning or action.

1 - With support: The message works after a prompt, clarification, or revision.

2 - Independently: The message is clear and accurate without a prompt.

Feedback sequence

First, identify a successful phrase. Next, identify one point where meaning or accuracy needs work. Offer one useful language correction, allow a repeat, and compare the two attempts. A total out of eight describes this performance only; do not translate it into a proficiency level.

Final challenge

Choose a case not rehearsed today. The learner gives a one-minute response, answers two follow-up questions, and writes a 70-110 word message. Add the lesson's second-round challenge. Compare the result with an earlier attempt using the four criteria.

Use the language responsibly

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Meaning before performance

Use specialist terms only when they help the listener. Explain unfamiliar words, preserve uncertainty, and ask when an instruction or responsibility is unclear. The model responses illustrate language; they do not authorize professional actions.

Sources and further reading

The cases, workshops, and explanations are original teaching material. These references provide language frameworks and selected professional context. References were checked in September 2026; consult current local guidance for actual work.

Digital.gov: plain-language guidance

<https://digital.gov/guides/plain-language>

Council of Europe: CEFR descriptors

<https://www.coe.int/en/web/common-european-framework-reference-languages/cefr-descriptors>

Your next step

Choose one expression you can use in a genuine work situation. Rehearse it with fictional details, then adapt the wording to your role and audience.