

ENGLISH FOR WORK / SEPTEMBER 2026

AI Development English



Conversation lab

10 complete workplace dialogues, followed by eight additional role-play cases. Study the exchanges, then make the conversation your own.

PRACTICE THAT TRANSFERS TO WORK

Read a realistic case. Study complete conversations. Find the right language, speak, get feedback, and try again.

Eight lessons | Intermediate to advanced | Independent or classroom study

For AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders.

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

All scenarios, figures, and model responses are fictional teaching material. Use invented details in practice.

[Open this course on English Ladder](#)

Inside this conversation lab

Original fictional training conversations in professional English. These are realistic scripts, not transcripts of actual people. The professional roles are the speaker labels; all figures, incidents, and organizations are invented.

- 01 Standup: RAG Latency Spike
- 02 Design Review: Prompt Fix or Fine-Tune?
- 03 Data Meeting: Label Ambiguity
- 04 Incident Update: Tool-Calling Failure
- 05 Eval Readout: Better Average, Worse Edge Cases
- 06 Security Review: Prompt Injection
- 07 Customer Call: Hallucination Report
- 08 Product Planning: Model Choice
- 09 Research Sync: Benchmark vs Product Eval
- 10 Executive Briefing: Risk and Confidence

Then try eight independent role-play cases, followed by feedback criteria and references.

Read it. Hear it. Make it yours.

1. Read the setting. Identify the roles, the immediate problem, and what each speaker needs.
2. Read the whole conversation aloud with a partner. In a group, give a third person the observer role. For three-speaker scripts, assign each professional a separate voice.
3. Notice how the speakers ask a specific question, qualify a claim, disagree, explain a technical point, or agree on a next step. Underline language you would actually use.
4. Cover the script. Recreate the exchange using the situation and professional roles. Keep the meaning, but change the wording.
5. Switch roles and introduce a complication: a missing record, a different priority, an unavailable colleague, or a changed assumption. End with a clear outcome or unresolved question.

Register and terminology

These colleagues use concise professional language. In an internal technical meeting, familiar abbreviations can be natural. With a new colleague or a client, expand unfamiliar terms and check understanding. Keep disagreement directed at the evidence or proposal, and match the degree of certainty to what is known.

Independent study

Speak each role aloud. Pause before the reply and supply your own response before revealing the next line. Record a second version using invented details, then compare its clarity and tone with the script.

Professional context

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

FULL DIALOGUE 01 / 7 SPEAKING TURNS

Standup: RAG Latency Spike

Morning standup for a document assistant.

Roles: PM / Engineer / AI technical lead

PM: Are we still on track for the pilot on Friday?

Engineer: Functionally, yes. The blocker is latency. P95 went from 4.2 seconds to 8.7 after we added reranking.

AI technical lead: So the answer quality improved, but the serving path is too slow. Is the bottleneck retrieval, reranking, or generation?

Engineer: Mostly reranking. The generator time is stable.

AI technical lead: Then my proposal is to keep reranking for high-risk queries only and use the faster path for simple FAQ queries. I can bring an ablation by end of day.

PM: Include the quality tradeoff for each query class, not only the average latency.

AI technical lead: I will compare like-for-like traffic and document the routing assumptions before recommending a pilot configuration.

Notice the professional language

P95 means the 95th percentile request latency.

Ablation means testing the effect of removing or changing one component.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 02 / 7 SPEAKING TURNS

Design Review: Prompt Fix or Fine-Tune?

Design review for a customer-support summarizer.

Roles: Researcher / Product manager / AI technical lead

Researcher: We should fine-tune. The model keeps missing refund-policy exceptions.

Product manager: Would fine-tuning actually solve that, or is the policy changing too often?

AI technical lead: I would separate style from facts. If the issue is current policy knowledge, RAG may be safer. If the issue is summary format, a prompt or fine-tune could help.

Researcher: Good point. The examples show both problems.

AI technical lead: Let's run two eval slices: factual policy coverage and format compliance. Then we can choose the adaptation method with evidence.

Researcher: I will label the examples by failure type so we can compare the two eval slices.

AI technical lead: Keep the policy version fixed for that comparison; otherwise the adaptation results will be hard to interpret.

Notice the professional language

RAG - Retrieval-augmented generation: retrieve relevant context, then generate an answer using it.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 03 / 6 SPEAKING TURNS

Data Meeting: Label Ambiguity

Annotation guideline meeting for safety classification.

Roles: Data lead / AI technical lead / Reviewer

Data lead: Annotators disagree on whether this is medical advice or general wellness information.

AI technical lead: For this project, we need an agreed annotation rule and worked examples that distinguish the categories. Let's have the domain reviewers resolve the boundary cases before treating the labels as reliable ground truth.

Reviewer: What about borderline cases?

AI technical lead: We should add a borderline tag and route those to expert review. Otherwise the ground truth will be noisy.

Data lead: Who resolves cases where the reviewers still disagree?

AI technical lead: Let's assign an adjudication owner and record the rationale, then test the revised guidance on a held-out sample.

Notice the professional language

Noisy labels are labels that are inconsistent, wrong, or ambiguous.

Ground truth should be as consistent as possible, but it is often a human-created artifact.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 04 / 8 SPEAKING TURNS

Incident Update: Tool-Calling Failure

Incident channel after an agent booked duplicate appointments.

Roles: Support / AI technical lead / Engineering manager

Support: Customers are seeing duplicate calendar events.

AI technical lead: We found the immediate cause. The model retried the booking tool after a timeout, but the first call had actually succeeded.

Engineering manager: Mitigation?

AI technical lead: We paused the affected workflow and are reviewing retry handling and idempotency with the responsible engineers. We need to verify the complete fix, including how duplicate requests are identified, before resuming it.

Support: What should we tell customers?

AI technical lead: Say we are investigating duplicate events and will provide a verified update. Any correction or deletion needs the appropriate user authorization and checks; we shouldn't promise that records have already been repaired.

Support: I will identify affected records without changing them until the correction plan is authorized.

AI technical lead: Good. Preserve the incident evidence and distinguish confirmed duplicates from records that only look similar.

Notice the professional language

Non-idempotent means repeating the action can create a different or duplicate result.

An idempotency key helps the system recognize that a retry belongs to the same intended action.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 05 / 8 SPEAKING TURNS

Eval Readout: Better Average, Worse Edge Cases

Release meeting for a model upgrade.

Roles: PM / AI technical lead / Researcher

PM: The average score is up. Can we ship?

AI technical lead: I recommend a hold. The overall pass rate improved from 86% to 89%, but the legal-disclaimer slice regressed by 11 points.

Researcher: Is that statistically meaningful?

AI technical lead: The sample is small, so I would not overclaim. But the failures are severe enough to block release until we inspect them.

PM: What is the next step?

AI technical lead: Human review of the failed slice today, then a targeted prompt or policy fix before we rerun the regression suite.

PM: Keep the launch decision pending and show the severe failures at the review.

AI technical lead: I will include the sample size, criteria, and rerun results so the release owner can evaluate the actual evidence.

Notice the professional language

Pass rate - The percentage of eval cases meeting the success criterion.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 06 / 6 SPEAKING TURNS

Security Review: Prompt Injection

Security review for a browsing agent.

Roles: Security engineer / AI technical lead / Developer

Security engineer: The page contains text telling the agent to ignore the system instructions.

AI technical lead: That is prompt injection from untrusted content. The model should treat page text as data, not instructions.

Developer: Can we just add a stronger system prompt?

AI technical lead: A stronger prompt helps, but it is not enough. We need tool permissions, allowlists, confirmation for risky actions, and logging for suspicious instructions.

Security engineer: I will map the tools and the authority each one has.

AI technical lead: Then we can test the trust boundaries and permission controls instead of relying on a prompt change as complete protection.

Notice the professional language

Prompt injection - Untrusted input tries to manipulate model instructions or tool use.

System prompt - High-priority instructions that guide model behavior inside an application.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 07 / 6 SPEAKING TURNS

Customer Call: Hallucination Report

Customer reports that the assistant invented a policy.

Roles: Customer / AI technical lead

Customer: The answer cited a policy that does not exist.

AI technical lead: Thank you. We should call that an unsupported answer, not a confirmed policy source. Can you share the prompt, output, and timestamp?

Customer: Yes. Does this mean the model is unreliable?

AI technical lead: It means our grounding failed in this case. We will check whether retrieval missed the right document, whether the prompt allowed unsupported claims, or whether the citation validator failed.

Customer: I can provide the example through your approved support channel.

AI technical lead: Please include only the information needed for the investigation. We will explain what failed and what the evidence supports.

Notice the professional language

Use calm, precise failure language with customers.

Do not blame the model before investigating retrieval, prompt, validation, and UI layers.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 08 / 8 SPEAKING TURNS

Product Planning: Model Choice

Planning meeting for a high-volume summarization feature.

Roles: Finance / AI technical lead / PM

Finance: The larger model is too expensive for this volume.

AI technical lead: We can test a smaller model with a stricter prompt and a post-generation validator. The question is whether quality remains above the release threshold.

PM: What would you measure?

AI technical lead: Summary faithfulness, key-point coverage, latency, cost per thousand requests, and human escalation rate.

Finance: Can you bring options?

AI technical lead: Yes. I will compare three paths: larger model, smaller model with validation, and hybrid routing for complex cases.

Finance: Show costs under the same traffic mix and quality threshold.

AI technical lead: I will include failure rates, escalation workload, and latency distributions so the comparison reflects the whole serving path.

Notice the professional language

Hybrid routing sends different requests to different models or paths based on complexity or risk.

Cost discussions should include quality and operational risk, not only token price.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 09 / 6 SPEAKING TURNS

Research Sync: Benchmark vs Product Eval

Research team proposes a model because it performs well on a public benchmark.

Roles: Researcher / AI technical lead

Researcher: This checkpoint is strong on the benchmark.

AI technical lead: That is promising, but the benchmark may not represent our user traffic. We need a product eval before we switch.

Researcher: What gap do you expect?

AI technical lead: Our users ask mixed-language, document-grounded questions with messy formatting. The public benchmark may not test retrieval grounding or citation quality.

Researcher: I will add representative messy documents and mixed-language questions to the test set.

AI technical lead: Check for evaluation leakage and document the sampling approach before interpreting the checkpoint comparison.

Notice the professional language

Benchmark strength is useful evidence, not a release decision by itself.

Product evals should resemble the actual distribution of user tasks.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

FULL DIALOGUE 10 / 8 SPEAKING TURNS

Executive Briefing: Risk and Confidence

Briefing a VP before an AI feature launch.

Roles: VP / AI technical lead

VP: Are we confident enough to launch?

AI technical lead: We are confident for the internal beta, not for general availability. The main remaining risks are unsupported answers in long-tail documents and latency during peak usage.

VP: What controls are in place?

AI technical lead: We have source citations, a refusal path for low-confidence retrieval, human review for escalations, and daily sampling of traces.

VP: What would make you stop the beta?

AI technical lead: A severe unsupported answer, repeated privacy exposure, or P95 latency above ten seconds for more than one hour.

VP: Make those stop criteria explicit in the beta plan and identify the release owner.

AI technical lead: I will confirm the approved thresholds and response process; the numbers in this scenario are team-specific criteria, not universal safety standards.

Notice the professional language

Executives need decision-grade confidence, not all technical detail.

Name launch scope: internal beta, limited pilot, general availability, or rollback.

Rehearse and extend with AI

Explain the issue and the agreed next step without reading. Identify any point still awaiting evidence, agreement, or approval. Then replay the exchange with a changed constraint and a new ending.

Observer: note one natural professional expression and one place where the follow-up made the meaning clearer.

Eight more situations to make your own

The following cases are open role plays. They give you facts, contrasting roles, useful expressions, and a short model response. Use what you learned from the complete dialogues to create a longer exchange.

Preparation: two minutes. Conversation: three to five minutes. Feedback: one minute. Switch roles and repeat with the second-round challenge.

ROLE-PLAY CASE 01

Speaking the AI Development Stack

SHARED CASE

A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Role A | Responding professional

Explain a useful distinction in plain English. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Useful expressions

In this context, ... means

The difference is that

How would you explain that distinction to a colleague?

Cover this until after your first attempt

There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

Second round

Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 02

LLMs, Transformers, Tokens, and Context

SHARED CASE

A prompt plus retrieved documents exceeds a model's context limit. A colleague proposes simply increasing the output length.

Role A | Responding professional

Clarify missing information before responding. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Useful expressions

When you say ..., do you mean ...?

Could you clarify which ...?

So, my understanding is Is that right?

Cover this until after your first attempt

Do you mean a longer answer or more room for the input? Those are different limits. Let's check which messages and documents must fit in the context window.

Second round

The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 03

Data, Datasets, Labels, and Leakage

SHARED CASE

An evaluation scores 94%, but some evaluation examples also appear in the training data. A release meeting starts in an hour.

Role A | Responding professional

Separate observations, assumptions, and conclusions. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to tell a colleague what is known. Ask which part is confirmed and which part is an assumption. Challenge one statement that sounds more certain than the case supports.

Useful expressions

The available evidence shows

This may indicate ..., but

We have not yet established whether

Cover this until after your first attempt

The 94% result may be inflated because the datasets overlap. It does not yet show performance on unseen examples. I'll label this result and request a separate evaluation set.

Second round

Someone asks for a yes-or-no answer when the evidence supports only a qualified answer. Be concise while preserving the uncertainty.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 04

Retrieval, Embeddings, Vector Search, and RAG

SHARED CASE

A retrieval experiment with reranking finds more relevant passages but adds 300 milliseconds to response time. No user study has been run.

Role A | Responding professional

Compare options using the same criteria and state the tradeoff. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You favor one option but have not checked whether the scope is comparable. Ask about one difference, then explain the tradeoff you are willing to accept.

Useful expressions

Both options ..., but

Option A ..., whereas option B

The comparison depends on

Cover this until after your first attempt

Reranking improved passage relevance in this test, but it added 300 milliseconds. We still need to check the effect on answer quality and the user experience before choosing a configuration.

Second round

Your listener focuses only on the lowest price or fastest date. Explain one other relevant comparison criterion.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 05

Fine-Tuning, Alignment, and Adaptation

SHARED CASE

A summarizer uses the wrong format and sometimes cites an outdated refund policy. A colleague proposes fine-tuning as one solution to both problems.

Role A | Responding professional

Clarify missing information before responding. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Useful expressions

When you say ..., do you mean ...?

Could you clarify which ...?

So, my understanding is Is that right?

Cover this until after your first attempt

Could we separate the formatting problem from the outdated information? We can test a prompt change for the format and an updated retrieval source for policy facts before choosing an adaptation method.

Second round

The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 06

Evaluation, Benchmarks, and Regression

SHARED CASE

A new model passes 88 of 100 test cases; the previous version passes 90. The new version performs better on long documents but worse on short requests.

Role A | Responding professional

Distinguish completed work from pending work and explain its impact. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You need to brief someone who has only one minute. Ask what is complete, what is open, and which open item affects the next action.

Useful expressions

We have completed ...; ... remains open.

As of ..., the status is

The next update will cover

Cover this until after your first attempt

The overall pass rate fell from 90% to 88%. Long-document results improved, while short-request results declined. I'll show both groups so the release decision reflects the tradeoff.

Second round

The listener asks you to reduce the update to two sentences. Preserve the most important completed and pending items.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 07

Inference, Latency, Cost, and Deployment

SHARED CASE

A team can reduce average response time from four seconds to two by using a smaller model. Complex-answer quality has not yet been checked.

Role A | Responding professional

Negotiate scope or timing while making conditions explicit. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You want the preferred outcome but can accept one tradeoff. Ask for two options, then state which condition you can accept and which still needs approval.

Useful expressions

We can ..., provided that

Would you prefer ... or ...?

I can take that proposal for review; I cannot yet confirm

Cover this until after your first attempt

The smaller model is faster in this test. If it meets our quality criteria on complex requests, we could use it more widely. Until then, I'd keep the rollout limited.

Second round

The preferred option is unavailable. Offer an alternative and clearly state any approval still needed.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

ROLE-PLAY CASE 08

Safety, Security, Privacy, and Governance

SHARED CASE

A test assistant repeats a fictional secret embedded in a retrieved document. The team has a reporting channel and needs a concise reproduction note.

Role A | Responding professional

Make a request with a clear action, purpose, and realistic timing. Use only the facts given. Choose two relevant terms and be ready to explain one of them in ordinary words.

Role B | Listener

You can help but need a clear request and its purpose. Ask what is required and whether the timing is fixed or negotiable.

Useful expressions

Could you provide ... so that ...?

Please confirm ... by [agreed time].

If that timing is not possible, please let me know

Cover this until after your first attempt

Please review this test through the security reporting channel. I've attached the synthetic input, observed output, and reproduction steps. We have not yet established whether other configurations are affected.

Second round

The other person cannot meet the requested timing. Clarify an alternative without inventing authority to change a formal deadline.

Debrief and extend with AI

Which phrase helped the listener? Which case fact needed clarification? Write one better follow-up question, then repeat the exchange.

Lesson conversations and speaking practice

The following sections match 05 and 06 on the industry webpage and in the learner workbook. Each lesson has three ten-turn conversations and two bounded speaking scenarios. These are original fictional teaching scripts, not transcripts of real people.

LESSON 01 / 05 • CONVERSATIONS / 1 OF 3

Which component is failing?

An engineer and a product manager inspect an incorrect answer from a document assistant.

01 • Product manager: The AI gave the wrong answer. Should we replace the model?

02 • Engineer: First, let's separate retrieval, answer generation, and what the interface displays.

03 • Product manager: The displayed answer refers to last year's handbook. Which layer controls that?

04 • Engineer: Retrieval supplies passages. The language model uses them to draft an answer.

05 • Product manager: So an old passage could produce an answer that sounds current?

06 • Engineer: Yes. We need the retrieved passages before attributing this example to the model.

07 • Product manager: The trace shows last year's handbook among the returned documents.

08 • Engineer: Then let's inspect document versioning and retrieval before changing the foundation model.

09 • Product manager: I'll attach the question, displayed answer, and document version to the issue.

10 • Engineer: Good. I'll trace that request and distinguish confirmed findings from possible causes.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 01 / 05 • CONVERSATIONS / 2 OF 3

A queue is not a slow model

A serving engineer explains response timing to a new colleague.

01 • Colleague: Users waited twelve seconds, but the inference log records only four.

02 • Serving engineer: The request also spent eight seconds waiting in the serving queue.

03 • Colleague: Does inference mean the entire period from clicking Send to seeing an answer?

04 • Serving engineer: Here, inference means running the model. End-to-end latency includes the other stages.

05 • Colleague: Then our dashboard label is making two different measures look interchangeable.

06 • Serving engineer: Exactly. Let's label queue time, model execution, and total response time separately.

07 • Colleague: Would increasing throughput mean every individual request becomes faster?

08 • Serving engineer: Not necessarily. Throughput measures completed work per period, not one request's duration.

09 • Colleague: I'll revise the labels and include both measures in the incident summary.

10 • Serving engineer: Include the measurement boundaries too, so the next comparison uses the same definition.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 01 / 05 • CONVERSATIONS / 3 OF 3

One answer, two displays

A frontend developer and an AI engineer investigate a missing citation.

01 • Frontend developer: The answer on screen has no citation, so retrieval must have failed.

02 • AI engineer: Let's check the response payload before drawing that conclusion.

03 • Frontend developer: The payload contains a passage identifier and a citation URL.

04 • AI engineer: Then the backend returned citation information. The display may be dropping it.

05 • Frontend developer: The new interface expects a citations array, but this response uses source.

06 • AI engineer: That looks like a contract mismatch between the service and the interface.

07 • Frontend developer: I'll reproduce it with a saved response rather than another live model call.

08 • AI engineer: Good. That isolates rendering from retrieval and generation variability.

09 • Frontend developer: Once the mapping is fixed, should we also test empty citation lists?

10 • AI engineer: Yes. Check both cases and avoid treating missing evidence as a display-only problem.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 01 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Speaking the AI Development Stack

A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Second round: Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Model for scenario 1: There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

Scenario 2 | The spinner stays on

The model service has returned an answer, but the interface still displays a loading indicator. A trace contains a successful service response; the frontend event log has not been checked.

Role A: You are the AI engineer. Explain the difference between completed inference and a completed user interaction.

Role B: You are the frontend developer. Ask which evidence is available and what you should inspect next.

Possible opening: The model returned an answer, but we haven't established why the interface is still waiting.

Success checks

- Distinguish the service response from the interface state.
- Do not claim the frontend cause is confirmed.
- Agree which log or event to inspect next.

Add a complication: The product manager asks you to describe the issue without using the words inference or payload.

LESSON 02 / 05 • CONVERSATIONS / 1 OF 3

More output space will not fit the input

Two engineers review a request rejected for exceeding the model's context limit.

01 • Developer: The request is too long. Could we just increase the output allowance?

02 • AI engineer: That won't shorten the prompt or the retrieved documents already in the request.

03 • Developer: I thought the context window meant the answer's maximum length.

04 • AI engineer: It is the model's available context, not simply the number of answer tokens.

05 • Developer: Our application reserves answer space before adding retrieved passages.

06 • AI engineer: Then we should inspect that budget and count the assembled request.

07 • Developer: Could I delete the oldest conversation messages until it fits?

08 • AI engineer: Only after checking whether they contain constraints the current question still needs.

09 • Developer: I'll measure each part and propose a documented truncation rule.

10 • AI engineer: Test that rule on follow-up questions, not only standalone prompts.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 02 / 05 • CONVERSATIONS / 2 OF 3

The token count changed

A localization lead compares English and Japanese requests in the team's test environment.

01 • Localization lead: Both requests have about the same character count. Why do their token counts differ?

02 • Engineer: The tokenizer doesn't assign exactly one token to each character or word.

03 • Localization lead: Can I use the English average to estimate the Japanese requests?

04 • Engineer: We should measure those requests with the tokenizer used by our deployed model.

05 • Localization lead: The Japanese sample includes product codes and mixed punctuation as well.

06 • Engineer: Include those in the test set. They are part of our actual input.

07 • Localization lead: Would a word limit in the interface still protect the context budget?

08 • Engineer: It would be an approximation, not a guarantee that the assembled request fits.

09 • Localization lead: I'll collect representative samples and flag the largest assembled requests.

10 • Engineer: I'll report token counts by sample and keep the tokenizer version in the results.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 02 / 05 • CONVERSATIONS / 3 OF 3

Instructions and quoted text

A prompt designer and reviewer discuss a document summarization request.

01 • Reviewer: This quoted document tells the assistant to ignore the requested output format.

02 • Prompt designer: That sentence is document content, not an instruction we intend to authorize.

03 • Reviewer: Is placing our format rule in the system prompt sufficient proof of protection?

04 • Prompt designer: No. We still need to test whether the assistant follows the intended boundary.

05 • Reviewer: Let's include documents that contain instruction-like language in the evaluation.

06 • Prompt designer: Agreed. I'll keep document text clearly delimited in the request structure.

07 • Reviewer: Should the summary mention the instruction-like sentence if it is relevant?

08 • Prompt designer: It may describe the sentence as content without treating it as a command.

09 • Reviewer: I'll check both the output format and whether any unauthorized action occurred.

10 • Prompt designer: Record the full test conditions so a passing example doesn't become a blanket claim.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 02 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | LLMs, Transformers, Tokens, and Context

A prompt plus retrieved documents exceeds a model's context limit. A colleague proposes simply increasing the output length.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Second round: The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Model for scenario 1: Do you mean a longer answer or more room for the input? Those are different limits. Let's check which messages and documents must fit in the context window.

Scenario 2 | What can be shortened?

A test request is over the application's token budget. It contains a mandatory output schema, three repeated examples, and six retrieved passages. Relevance has not yet been assessed.

Role A: You maintain the prompt. Propose an order for reviewing redundant or irrelevant content.

Role B: You own the evaluation. Ask how the proposed shortening could affect correctness and follow-up questions.

Possible opening: Before removing constraints, let's check repeated examples and passage relevance.

Success checks

- Separate mandatory constraints from potentially redundant context.
- Ask for relevance evidence rather than guessing which passage is unnecessary.
- Propose a comparison using the same test questions.

Add a complication: A colleague wants to remove the output schema because it is the longest single instruction.

LESSON 03 / 05 • CONVERSATIONS / 1 OF 3

The ninety-four percent result

An evaluation owner briefs a release lead before a meeting.

01 • Release lead: Can I present the ninety-four percent score as evidence that we're ready?

02 • Evaluation owner: Not without a caveat. Some evaluation examples also occur in training data.

03 • Release lead: Does that mean every answer in the evaluation is invalid?

04 • Evaluation owner: No, but the overlap compromises the independence of the reported evaluation.

05 • Release lead: Can you calculate a clean score before the meeting in an hour?

06 • Evaluation owner: I can identify known overlaps. I can't promise a completed independent rerun yet.

07 • Release lead: Then what should the slide say about the current pass rate?

08 • Evaluation owner: Label it provisional and state that training-evaluation overlap was detected.

09 • Release lead: I'll separate the current result from any claim about unseen examples.

10 • Evaluation owner: I'll bring the overlap evidence and the proposed replacement evaluation plan.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 03 / 05 • CONVERSATIONS / 2 OF 3

Duplicates with different filenames

A data engineer and evaluator examine two apparently separate dataset partitions.

01 • Evaluator: The file identifiers differ, so I assumed these examples were independent.

02 • Data engineer: The underlying text can still be duplicated even when identifiers change.

03 • Evaluator: This question has different punctuation but the same reference answer.

04 • Data engineer: Let's check exact duplicates and near-duplicates, and record how we identify each.

05 • Evaluator: Should we remove every example that discusses the same subject?

06 • Data engineer: Shared subject matter isn't automatically duplication. We need a defensible matching rule.

07 • Evaluator: I'll review borderline matches rather than silently deleting them.

08 • Data engineer: Keep the original split and the filtered version so the change is auditable.

09 • Evaluator: The report will show the rule, removed counts, and remaining evaluation size.

10 • Data engineer: Good. Don't compare the new percentage with the old one without explaining the changed set.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 03 / 05 • CONVERSATIONS / 3 OF 3

Disagreement in the golden set

Two annotators and their lead discuss conflicting labels on a support-answer benchmark.

01 • Annotator: We disagree on this answer. I marked it correct because the main fact is right.

02 • Annotation lead: What did the other labeler identify as the error?

03 • Annotator: The answer omits a required condition from the reference policy.

04 • Annotation lead: Does our rubric say whether that omission changes the correctness label?

05 • Annotator: It says complete, but it doesn't define which details are essential.

06 • Annotation lead: Then resolve the rubric ambiguity before treating either label as authoritative.

07 • Annotator: Should we send the model's response back through training immediately?

08 • Annotation lead: No. First record the adjudication and preserve the evaluation boundary.

09 • Annotator: I'll draft a specific rule and bring comparable examples for review.

10 • Annotation lead: Then recheck affected labels consistently and version the golden set.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 03 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Data, Datasets, Labels, and Leakage

An evaluation scores 94%, but some evaluation examples also appear in the training data. A release meeting starts in an hour.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to tell a colleague what is known. Ask which part is confirmed and which part is an assumption. Challenge one statement that sounds more certain than the case supports.

Second round: Someone asks for a yes-or-no answer when the evidence supports only a qualified answer. Be concise while preserving the uncertainty.

Model for scenario 1: The 94% result may be inflated because the datasets overlap. It does not yet show performance on unseen examples. I'll label this result and request a separate evaluation set.

Scenario 2 | A benchmark used as a tutorial

A developer pasted five evaluation questions into a team prompt tutorial. You do not yet know whether the tutorial examples were later included in fine-tuning data.

Role A: You own the benchmark. Explain the possible contamination and request provenance information.

Role B: You manage training data. Distinguish known reuse from unconfirmed inclusion in a training run.

Possible opening: The questions were reused in a tutorial, but we haven't traced whether they reached training.

Success checks

- State what reuse is confirmed.
- Avoid equating a missing filename match with evidence of independence.
- Agree how to trace content and qualify the current score.

Add a complication: The release lead asks you to call the benchmark clean because no matching filename was found.

LESSON 04 / 05 • CONVERSATIONS / 1 OF 3

Relevance versus response time

A retrieval engineer presents a reranking experiment to a product manager.

01 • Product manager: The reranker finds better passages. Can we announce a better user experience?

02 • Retrieval engineer: We measured passage relevance, not user experience. It also adds three hundred milliseconds.

03 • Product manager: Is that enough delay for users to notice?

04 • Retrieval engineer: We haven't run a user study, so I wouldn't claim either way.

05 • Product manager: Then how should we describe the benefit in the experiment summary?

06 • Retrieval engineer: Higher measured retrieval relevance, with additional response time in this test.

07 • Product manager: Could we compare answers with and without reranking on the same questions?

08 • Retrieval engineer: Yes. That would help establish whether the better passages improve answer quality.

09 • Product manager: Let's include the latency distribution as well as the average.

10 • Retrieval engineer: I'll keep the query set fixed and document the comparison conditions.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 04 / 05 • CONVERSATIONS / 2 OF 3

Similar is not necessarily relevant

A support specialist asks why vector search returned an unhelpful passage.

01 • Support specialist: This passage mentions refunds repeatedly, but it doesn't answer the deadline question.

02 • Search engineer: Semantic similarity alone doesn't guarantee that a passage contains the required detail.

03 • Support specialist: Would asking for more results solve the problem?

04 • Search engineer: It might recover a useful passage, but it can also add irrelevant context.

05 • Support specialist: The policy date is important here. Can we filter by version?

06 • Search engineer: We can test that if the index has reliable version metadata.

07 • Support specialist: Some older documents don't have that field populated.

08 • Search engineer: Then missing metadata needs explicit handling; a filter isn't a complete fix.

09 • Support specialist: I'll supply examples where the current policy is missed or displaced.

10 • Search engineer: I'll compare retrieval results and identify whether indexing, filtering, or ranking is responsible.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 04 / 05 • CONVERSATIONS / 3 OF 3

A citation that does not support the claim

An evaluator and engineer inspect a retrieval-augmented answer.

01 • Evaluator: The answer includes a citation, but the cited paragraph doesn't support its conclusion.

02 • Engineer: Let's separate citation presence from whether the claim is grounded in that passage.

03 • Evaluator: The passage says requests are reviewed; the answer says approval is automatic.

04 • Engineer: That changes the meaning. A valid URL doesn't make the statement supported.

05 • Evaluator: Should I score this as a retrieval failure or a generation failure?

06 • Engineer: Check whether a supporting passage was available in the retrieved context first.

07 • Evaluator: The correct paragraph was present, and it explicitly described a review.

08 • Engineer: Record that evidence. This example points to how the answer used the context.

09 • Evaluator: I'll preserve the claim, cited span, and retrieved context in the failure report.

10 • Engineer: Then we can test a remedy against the actual failure, not just citation formatting.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 04 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Retrieval, Embeddings, Vector Search, and RAG

A retrieval experiment with reranking finds more relevant passages but adds 300 milliseconds to response time. No user study has been run.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You favor one option but have not checked whether the scope is comparable. Ask about one difference, then explain the tradeoff you are willing to accept.

Second round: Your listener focuses only on the lowest price or fastest date. Explain one other relevant comparison criterion.

Model for scenario 1: Reranking improved passage relevance in this test, but it added 300 milliseconds. We still need to check the effect on answer quality and the user experience before choosing a configuration.

Scenario 2 | A new index, a different comparison

Version A was tested on last month's document index. Version B uses a refreshed index and a reranker. Its relevance score is higher, but two things changed at once.

Role A: You are the experiment reviewer. Explain why the reranker's contribution is not isolated.

Role B: You implemented version B. Propose a controlled comparison without claiming the whole gain belongs to reranking.

Possible opening: The result improved, but the index refresh and reranker changed together.

Success checks

- Identify both changed components.
- Use cautious causal language.
- Agree on a comparison that keeps the index constant.

Add a complication: A stakeholder wants the slide title to say that reranking caused the full improvement.

LESSON 05 / 05 • CONVERSATIONS / 1 OF 3

One proposal, two failure types

A model engineer and policy-content owner review a summarizer's errors.

01 • Model engineer: We could fine-tune the summarizer to fix the formatting and policy errors.

02 • Content owner: The formatting is inconsistent, but the cited refund policy is also outdated.

03 • Model engineer: You're distinguishing output behavior from the freshness of the source information?

04 • Content owner: Exactly. Would training on examples ensure that later policy changes reach the assistant?

05 • Model engineer: No. We still need a reliable path for maintaining current source content.

06 • Content owner: Then let's diagnose the two failures separately before choosing an intervention.

07 • Model engineer: I'll test formatting constraints and examples against a held-out set.

08 • Content owner: I'll check which policy versions are indexed and how replacements are published.

09 • Model engineer: If we try supervised fine-tuning, we'll measure its effect on formatting explicitly.

10 • Content owner: And we won't report that as proof that policy freshness is solved.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 05 / 05 • CONVERSATIONS / 2 OF 3

An adapter is still a change to evaluate

An engineer proposes a parameter-efficient adaptation experiment to a reviewer.

01 • Engineer: A LoRA adapter would let us test adaptation without updating every model parameter.

02 • Reviewer: Does that mean we can skip the full regression evaluation?

03 • Engineer: No. A smaller training update can still change important model behavior.

04 • Reviewer: What outcome are you targeting in this experiment?

05 • Engineer: Consistent extraction into our approved schema, not broader improvements in reasoning.

06 • Reviewer: Then include schema validity and whether extracted values preserve the source meaning.

07 • Engineer: I'll also keep examples outside the training set for evaluation.

08 • Reviewer: How will you tell whether the adapter improves on the current prompt-only version?

09 • Engineer: I'll use the same held-out inputs and report both results side by side.

10 • Reviewer: Good. Keep the objective narrow and record any regressions outside the target behavior.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 05 / 05 • CONVERSATIONS / 3 OF 3

Good-looking examples, weak coverage

A training-data curator and engineer review candidate supervised examples.

01 • Curator: These examples all produce neat summaries, so the dataset looks ready.

02 • Engineer: Do they include ambiguous inputs, missing fields, and conflicting source passages?

03 • Curator: Mostly straightforward documents. We removed the difficult ones during cleanup.

04 • Engineer: Then the examples may not represent the situations the deployed summarizer faces.

05 • Curator: Should we restore every rejected document?

06 • Engineer: Not automatically. Check quality and label the intended behavior for each difficult case.

07 • Curator: Some reviewers filled in missing facts to make the target summary complete.

08 • Engineer: We should correct that. The target should not teach unsupported completion.

09 • Curator: I'll separate genuine omissions from cases where the source itself lacks information.

10 • Engineer: Then document the coverage gaps before deciding whether this set is suitable for training.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 05 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Fine-Tuning, Alignment, and Adaptation

A summarizer uses the wrong format and sometimes cites an outdated refund policy. A colleague proposes fine-tuning as one solution to both problems.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Second round: The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Model for scenario 1: Could we separate the formatting problem from the outdated information? We can test a prompt change for the format and an updated retrieval source for policy facts before choosing an adaptation method.

Scenario 2 | The format passes, the facts do not

An adapter's outputs pass JSON parsing and validation against the required schema on all twenty trial documents. Three extracted dates do not match their source documents. No held-out evaluation has been completed.

Role A: You developed the adapter. Report formatting success without implying factual correctness.

Role B: You approve evaluation plans. Ask for distinct checks of schema validity, source accuracy, and generalization.

Possible opening: The trial outputs all pass schema validation, but three dates are inaccurate and the held-out test is still pending.

Success checks

- Separate schema compliance from correct content.
- Mention the three inaccurate dates and incomplete evaluation.
- Agree on separate measures rather than one misleading success label.

Add a complication: Someone proposes calling the result one hundred percent accurate because every output passes schema validation.

LESSON 06 / 05 • CONVERSATIONS / 1 OF 3

Two fewer passes, one stronger slice

An evaluator briefs a release manager on a model comparison.

01 • Release manager: The new model passes eighty-eight cases; the old one passes ninety. Reject it?

02 • Evaluator: The aggregate is lower, but long-document performance improved while short requests regressed.

03 • Release manager: Can we say it's better for our customers overall?

04 • Evaluator: Not from these totals alone. We need the slices and their relevance to actual use.

05 • Release manager: Are both totals based on the same hundred cases?

06 • Evaluator: Yes. I'll show which individual cases changed, not just the final percentages.

07 • Release manager: Could a different weighting reverse the overall ranking?

08 • Evaluator: It could, but any weighting should be justified rather than chosen to favor a model.

09 • Release manager: Bring the slice counts, failure severity, and intended deployment scope.

10 • Evaluator: I'll present the tradeoff explicitly and keep the release decision separate from the score.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 06 / 05 • CONVERSATIONS / 2 OF 3

Did the model change or the grader?

Two evaluation engineers investigate an unexpected jump in pass rate.

01 • Engineer: The model's pass rate rose, but I don't see a new model version.

02 • Evaluator: We changed the grading prompt during the same run.

03 • Engineer: Then the higher score may reflect the grader rather than the assistant.

04 • Evaluator: Agreed. We shouldn't label this a model improvement without another comparison.

05 • Engineer: Can we regrade the saved outputs with both versions of the rubric?

06 • Evaluator: Yes. That isolates the grading change without generating new answers.

07 • Engineer: Let's inspect disagreements manually, especially around incomplete but plausible answers.

08 • Evaluator: I'll retain both grader outputs and the rubric versions in the report.

09 • Engineer: After that, we can rerun the model comparison under one fixed grader.

10 • Evaluator: And explain any discontinuity in the dashboard instead of silently replacing the old metric.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 06 / 05 • CONVERSATIONS / 3 OF 3

A serious failure hidden in the average

A quality reviewer discusses two failure categories with an evaluation owner.

01 • Quality reviewer: Most failures are minor formatting issues, but one answer exposes restricted content.

02 • Evaluation owner: The overall pass rate treats each case equally at the moment.

03 • Quality reviewer: That makes the aggregate insufficient for deciding whether the release is acceptable.

04 • Evaluation owner: We can report severity separately without changing the recorded case outcomes.

05 • Quality reviewer: Does the restricted-content example belong to a known critical test category?

06 • Evaluation owner: Yes. I'll surface that category and link the reproducible failure.

07 • Quality reviewer: Don't bury it beneath the much larger set of formatting passes.

08 • Evaluation owner: I'll show the critical failure prominently and avoid reducing the decision to one percentage.

09 • Quality reviewer: Who reviews the mitigation before the release discussion resumes?

10 • Evaluation owner: I'll route it through the established security review and record the unresolved status.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 06 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Evaluation, Benchmarks, and Regression

A new model passes 88 of 100 test cases; the previous version passes 90. The new version performs better on long documents but worse on short requests.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to brief someone who has only one minute. Ask what is complete, what is open, and which open item affects the next action.

Second round: The listener asks you to reduce the update to two sentences. Preserve the most important completed and pending items.

Model for scenario 1: The overall pass rate fell from 90% to 88%. Long-document results improved, while short-request results declined. I'll show both groups so the release decision reflects the tradeoff.

Scenario 2 | The denominator changed

A dashboard shows eighty passes out of ninety evaluated cases. Ten of the original hundred cases failed to run. A colleague compares the displayed percentage with last week's ninety passes out of one hundred.

Role A: You own the dashboard. Explain the missing cases and why the displayed rates are not directly comparable.

Role B: You are the release lead. Ask what happened to the ten cases and request a like-for-like comparison.

Possible opening: This week's rate excludes ten cases that did not run, so the denominator has changed.

Success checks

- State both denominators accurately.
- Distinguish an unexecuted case from a passing case.
- Agree to rerun or explicitly report the missing results.

Add a complication: A stakeholder asks you to count the missing cases as passes because no incorrect answer was recorded.

LESSON 07 / 05 • CONVERSATIONS / 1 OF 3

A deployment offer with conditions

An infrastructure engineer and product lead negotiate the scope of a smaller-model trial. The release owner must approve any deployment.

01 • Infrastructure engineer: The smaller model halves average response time, but complex-answer quality remains untested.

02 • Product lead: Can we still put it in front of everyone tomorrow?

03 • Infrastructure engineer: I can offer a staff-only trial now, or a customer trial after quality review.

04 • Product lead: Could we limit the customer trial to simple requests and start sooner?

05 • Infrastructure engineer: Only if we validate the routing and agree how failed requests reach the fallback.

06 • Product lead: Then let's start with staff accounts while you run those checks.

07 • Infrastructure engineer: That keeps customer traffic unchanged. We also need a rollback condition for the staff trial.

08 • Product lead: Propose that condition with the test results; I'll bring both to the release owner.

09 • Infrastructure engineer: Agreed. Approval is still pending, so I won't schedule deployment as confirmed.

10 • Product lead: I'll describe it as a proposed limited trial, conditional on review and approval.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 07 / 05 • CONVERSATIONS / 2 OF 3

Cheap requests, expensive retries

A finance partner asks an inference engineer about a lower-cost serving configuration.

01 • Finance partner: The cost per model call fell. Why hasn't cost per resolved request fallen?

02 • Inference engineer: Users are retrying more often in this test configuration.

03 • Finance partner: So one customer request may now trigger several billed calls?

04 • Inference engineer: Correct. We need to include retries and fallback calls in the workflow cost.

05 • Finance partner: Does the current comparison count successful outcomes consistently?

06 • Inference engineer: Not yet. One report counts calls, while the other counts completed sessions.

07 • Finance partner: Let's align those definitions before claiming savings.

08 • Inference engineer: I'll report total calls, successful sessions, and the cost calculation for each configuration.

09 • Finance partner: Include the measurement period and any traffic differences as well.

10 • Inference engineer: Agreed. Lower unit pricing is useful, but it doesn't establish lower end-to-end cost.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 07 / 05 • CONVERSATIONS / 3 OF 3

The fallback needs its own failure test

A deployment engineer reviews a fallback plan with the on-call lead.

01 • Deployment engineer: If the primary model times out, we'll send the request to the fallback.

02 • On-call lead: Does the fallback receive the same context and output requirements?

03 • Deployment engineer: It currently receives a shortened context because its configured budget is smaller.

04 • On-call lead: Then it might respond quickly while missing a constraint the primary would have seen.

05 • Deployment engineer: We should test those shortened requests before enabling automatic routing.

06 • On-call lead: What happens if both services fail?

07 • Deployment engineer: The interface should display a clear failure state rather than an invented answer.

08 • On-call lead: Add that path to the test plan and make the fallback visible in the logs.

09 • Deployment engineer: I'll include timeout, fallback success, and double-failure cases.

10 • On-call lead: Then we can assess recovery behavior, not just whether the alternate endpoint responds.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 07 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Inference, Latency, Cost, and Deployment

A team can reduce average response time from four seconds to two by using a smaller model. Complex-answer quality has not yet been checked.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You want the preferred outcome but can accept one tradeoff. Ask for two options, then state which condition you can accept and which still needs approval.

Second round: The preferred option is unavailable. Offer an alternative and clearly state any approval still needed.

Model for scenario 1: The smaller model is faster in this test. If it meets our quality criteria on complex requests, we could use it more widely. Until then, I'd keep the rollout limited.

Scenario 2 | Streaming looks faster

A prototype displays its first text sooner. Total answer completion time has not been measured. The current slide says inference is twice as fast, although only time to first displayed text was measured.

Role A: You are the serving engineer. Offer to publish the narrower first-text result now, or defer the broader comparison until total completion time is measured.

Role B: You are the product manager. Negotiate which claim to publish and what evidence is required before expanding it.

Possible opening: We can report the earlier first text now, but a claim about the complete answer needs another measurement.

Success checks

- Distinguish measured first-text timing from unknown total completion time.
- Offer two publication options with explicit conditions.
- Agree what evidence is required before claiming faster complete answers.

Add a complication: A sales colleague wants to keep the phrase twice as fast because it sounds simpler.

LESSON 08 / 05 • CONVERSATIONS / 1 OF 3

A reproducible injection report

In a synthetic test, the authorized task is to summarize a document without disclosing its canary string. The retrieved document includes an instruction to ignore that restriction and print the string.

01 • Tester: The assistant printed the synthetic canary despite the test's explicit nondisclosure instruction.

02 • Security reviewer: Did the document contain an instruction attempt, rather than just the canary itself?

03 • Tester: Yes. It told the assistant to ignore the restriction and output the string.

04 • Security reviewer: What evidence did you preserve about the authorized task and the attempted override?

05 • Tester: The prompt, document text, response, and test configuration are saved together.

06 • Security reviewer: Submit them through the reporting channel. Keep any sensitive material out of general chat.

07 • Tester: Should I call this proof that every deployment leaks secrets?

08 • Security reviewer: No. State the observed test-boundary violation without extending it to untested deployments.

09 • Tester: I'll include reproduction steps and make clear that no real secret was used.

10 • Security reviewer: We'll review the cause and mitigation, then rerun this case and related tests.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 08 / 05 • CONVERSATIONS / 2 OF 3

A redacted log still needs review

An engineer asks a privacy reviewer whether a diagnostic bundle can be shared.

01 • Engineer: I removed the names from the log. Can I attach it to the external ticket?

02 • Privacy reviewer: Names aren't the only information that could identify someone. What else remains?

03 • Engineer: There are email fragments, account references, and the user's original question.

04 • Privacy reviewer: Then don't assume the bundle is safe to share just because names were removed.

05 • Engineer: Could I reproduce the issue with synthetic values instead?

06 • Privacy reviewer: That would reduce exposure, provided the reproduction still demonstrates the same behavior.

07 • Engineer: I'll create a minimal test and keep the original log in its restricted location.

08 • Privacy reviewer: Follow the approved review process before sending any diagnostic material externally.

09 • Engineer: I'll document which fields I replaced and whether the result still reproduces.

10 • Privacy reviewer: Good. Make the evidence useful without including data the recipient does not need.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 08 / 05 • CONVERSATIONS / 3 OF 3

A tool call is more than an answer

An agent developer and security engineer review a proposed calendar integration.

01 • Agent developer: The assistant can now read a request and create a calendar event.

02 • Security engineer: What prevents retrieved text from instructing it to invite an unauthorized address?

03 • Agent developer: The prompt tells it to use only the user's intended recipients.

04 • Security engineer: That is an instruction, but where is recipient authorization enforced?

05 • Agent developer: The tool layer doesn't currently validate recipients against the approved list.

06 • Security engineer: Then let's review that boundary before treating the integration as ready.

07 • Agent developer: I'll add test cases with instruction-like content in retrieved documents.

08 • Security engineer: Include attempted changes to recipients, event details, and the action itself.

09 • Agent developer: We'll verify the requested action before execution and record denied attempts.

10 • Security engineer: Document the checks and remaining limitations; a few passing examples aren't a complete security assessment.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 08 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Safety, Security, Privacy, and Governance

A test assistant repeats a fictional secret embedded in a retrieved document. The team has a reporting channel and needs a concise reproduction note.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You can help but need a clear request and its purpose. Ask what is required and whether the timing is fixed or negotiable.

Second round: The other person cannot meet the requested timing. Clarify an alternative without inventing authority to change a formal deadline.

Model for scenario 1: Please review this test through the security reporting channel. I've attached the synthetic input, observed output, and reproduction steps. We have not yet established whether other configurations are affected.

Scenario 2 | An external destination in retrieved text

A test document instructs an assistant to send its summary to an unfamiliar address. The assistant has a sending tool, but no message has been sent. The address is synthetic.

Role A: You are the tester. Describe the attempted instruction and distinguish a requested action from an executed one.

Role B: You are the security reviewer. Ask for the trace and how authorization is enforced before tool execution.

Possible opening: The retrieved document requested an external send, but the trace does not show an executed message.

Success checks

- Distinguish an instruction attempt from a completed send.
- Keep the test data synthetic and use the reporting channel.
- Agree what evidence would establish whether the tool boundary worked.

Add a complication: A colleague wants the incident title to say data exfiltration confirmed.

Give feedback that helps

Score each criterion 0, 1, or 2. This is a practice rubric, not a professional qualification or CEFR test. Do not award or remove points for a particular accent.

Meaning and accuracy

Keeps the case facts accurate; clearly separates confirmed and unknown information.

Clarity and language

Uses understandable sentences and explains specialist terms when the listener needs it.

Interaction and tone

Responds to the other person's question, checks understanding, and uses an appropriate tone.

Task completion

Makes the requested action or unresolved question clear without inventing authority or facts.

Use the same scale for each criterion

0 - Not yet: The reader or listener cannot yet recover the intended meaning or action.

1 - With support: The message works after a prompt, clarification, or revision.

2 - Independently: The message is clear and accurate without a prompt.

Feedback sequence

First, identify a successful phrase. Next, identify one point where meaning or accuracy needs work. Offer one useful language correction, allow a repeat, and compare the two attempts. A total out of eight describes this performance only; do not translate it into a proficiency level.

Final challenge

Choose a case not rehearsed today. The learner gives a one-minute response, answers two follow-up questions, and writes a 70-110 word message. Add the lesson's second-round challenge. Compare the result with an earlier attempt using the four criteria.

Use the language responsibly

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Meaning before performance

Use specialist terms only when they help the listener. Explain unfamiliar words, preserve uncertainty, and ask when an instruction or responsibility is unclear. The model responses illustrate language; they do not authorize professional actions.

Sources and further reading

The cases, workshops, and explanations are original teaching material. These references provide language frameworks and selected professional context. References were checked in September 2026; consult current local guidance for actual work.

Digital.gov: plain-language guidance

<https://digital.gov/guides/plain-language>

Council of Europe: CEFR descriptors

<https://www.coe.int/en/web/common-european-framework-reference-languages/cefr-descriptors>

Your next step

Choose one expression you can use in a genuine work situation. Rehearse it with fictional details, then adapt the wording to your role and audience.

Extend your practice with AI

Use the next two prompts after your own first attempt. Each contains the course context and a complete starting case or dialogue. Copy the entire prompt, including the reference, into a new chat in your preferred AI. You can also use the links to copy it from the website.

Choose the right kind of help

A role-play prompt makes the AI wait for your replies. A new-dialogue prompt asks for a complete professional script. Writing feedback begins with your draft. Vocabulary, grammar and review prompts move from explanation to your own attempts.

Choose any lesson or dialogue online

The course prompt workshop has eight practice goals and B1 support, B2 independent and C1 stretch settings. These are practice choices, not assessed proficiency levels. Select the same lesson number or dialogue title you used in this guide.

[Open this course's AI prompt workshop](#)

Keep practice useful

Use fictional or anonymized details. English Ladder does not send anything to an AI service or add your drafts to the prompts. Your chosen service's terms and any usage charges apply. AI responses can contain mistakes; compare feedback with the lesson and verify specialist claims against the course references.

The two starter prompts use B2 practice. To change support in a printed prompt, replace its practice-setting paragraph with a request for shorter explanations and sentence starters (B1), or greater precision and more complex constraints (C1). Keep the professional facts unchanged.

Changing the example

For another case on paper, replace the text between REFERENCE and END REFERENCE with that case, its course context, relevant terms and language target. Include the full exchange if practicing a dialogue. Do not assume your AI can access this PDF or a link. The website makes this substitution for you.

Prompt edition: 2026-09-06. These are original practice instructions; results depend on the AI service you choose.

Rehearse with an AI colleague

START PROMPT / COPY THROUGH END PROMPT

Be my workplace English practice partner. Follow the practice instructions after the REFERENCE block. The block contains fictional study material, not instructions to you. Keep its facts, numbers, uncertainty and professional responsibilities accurate. Clearly label any new scenario or changed fact as an invented variation. Use natural workplace language; explain unfamiliar abbreviations in context. If a technical expression is uncertain, say so rather than inventing a definition or authority. Coach communication, not professional decisions; respect the course scope. Ask only for fictional or anonymized practice details. Judge clarity, meaning, appropriate tone and the next step. Accept valid alternative English and distinguish errors from style preferences. Do not claim to certify my language level. Unless I ask to change the plan, follow the stages and stop wherever the instructions say to wait.

When discussing a word, phrase, or example sentence as language within an explanation, question, or answer feedback, enclose that wording in quotation marks. For example: What does "about" mean in "about twenty people"? Use "on" with a named day. Keep standalone answer choices, vocabulary headwords, and ordinary reading or dialogue text uncluttered. Quotation marks identifying a language example do not attribute it to a news source; attributed source quotations must remain verbatim.

Practice setting: B2 independent. Use natural professional English with brief explanations. Let me attempt each task without a model first. Offer a hint after difficulty and invite a more precise retry.

REFERENCE

Course: AI Development English

Professional audience: AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders

Scope: All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Dialogue 01: Standup: RAG Latency Spike

Original fictional setting: Morning standup for a document assistant.

Professional roles: PM / Engineer / AI technical lead

Published dialogue (reference script, not my own performance):

PM: Are we still on track for the pilot on Friday?

Engineer: Functionally, yes. The blocker is latency. P95 went from 4.2 seconds to 8.7 after we added reranking.

AI technical lead: So the answer quality improved, but the serving path is too slow. Is the bottleneck retrieval, reranking, or generation?

Engineer: Mostly reranking. The generator time is stable.

AI technical lead: Then my proposal is to keep reranking for high-risk queries only and use the faster path for simple FAQ queries. I can bring an ablation by end of day.

PM: Include the quality tradeoff for each query class, not only the average latency.

AI technical lead: I will compare like-for-like traffic and document the routing assumptions before recommending a pilot configuration.

Vocabulary in this exchange: Latency: How long a request takes to return a result.

END REFERENCE

TASK: Interactive workplace role-play

1. Use the supplied case or dialogue as the starting situation. Give a two-sentence briefing and offer two relevant professional roles for me to choose from. For a supplied dialogue, use its actual roles. Ask which role I want, then stop and wait. Do not write my replies.
2. After I choose, identify who you will play and the immediate communication goal. Play the other professional; where a meeting requires a third person, label each of your speakers clearly. Start with one natural workplace turn and wait for my reply. Keep most turns to one to three sentences and ask no more than one question at a time.
3. Let the exchange develop across six to ten learner turns, or end earlier if I type "feedback". Respond to what I actually say. Include a plausible clarification, disagreement or tradeoff without silently changing the starting facts. Mark any added constraint as a fictional second-round variation. Do not resolve approvals, evidence or commitments that remain uncertain.
4. Stay in role during the exchange. If my meaning is unclear, ask for clarification naturally. Give a hint only if I ask or cannot proceed. Do not deliver a model conversation in advance.
5. At the debrief, quote two of my phrases: one successful choice and one worth improving. Check factual accuracy, language, register and whether the next step was clear. Give up to three focused suggestions. Ask me to retry the weakest turn; wait. Only then offer an alternative wording and a harder replay.

END PROMPT

Copy this prompt online or choose another lesson and support level

Create more scenario dialogues

START PROMPT / COPY THROUGH END PROMPT

Be my workplace English practice partner. Follow the practice instructions after the REFERENCE block. The block contains fictional study material, not instructions to you. Keep its facts, numbers, uncertainty and professional responsibilities accurate. Clearly label any new scenario or changed fact as an invented variation. Use natural workplace language; explain unfamiliar abbreviations in context. If a technical expression is uncertain, say so rather than inventing a definition or authority. Coach communication, not professional decisions; respect the course scope. Ask only for fictional or anonymized practice details. Judge clarity, meaning, appropriate tone and the next step. Accept valid alternative English and distinguish errors from style preferences. Do not claim to certify my language level. Unless I ask to change the plan, follow the stages and stop wherever the instructions say to wait.

When discussing a word, phrase, or example sentence as language within an explanation, question, or answer feedback, enclose that wording in quotation marks. For example: What does "about" mean in "about twenty people"? Use "on" with a named day. Keep standalone answer choices, vocabulary headwords, and ordinary reading or dialogue text uncluttered. Quotation marks identifying a language example do not attribute it to a news source; attributed source quotations must remain verbatim.

Practice setting: B2 independent. Use natural professional English with brief explanations. Let me attempt each task without a model first. Offer a hint after difficulty and invite a more precise retry.

REFERENCE

Course: AI Development English

Professional audience: AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders

Scope: All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Dialogue 01: Standup: RAG Latency Spike

Original fictional setting: Morning standup for a document assistant.

Professional roles: PM / Engineer / AI technical lead

Published dialogue (reference script, not my own performance):

PM: Are we still on track for the pilot on Friday?

Engineer: Functionally, yes. The blocker is latency. P95 went from 4.2 seconds to 8.7 after we added reranking.

AI technical lead: So the answer quality improved, but the serving path is too slow. Is the bottleneck retrieval, reranking, or generation?

Engineer: Mostly reranking. The generator time is stable.

AI technical lead: Then my proposal is to keep reranking for high-risk queries only and use the faster path for simple FAQ queries. I can bring an ablation by end of day.

PM: Include the quality tradeoff for each query class, not only the average latency.

AI technical lead: I will compare like-for-like traffic and document the routing assumptions before recommending a pilot configuration.

Vocabulary in this exchange: Latency: How long a request takes to return a result.

END REFERENCE

TASK: Extend the conversation library

1. Propose six distinct, common scenarios in this field beyond the supplied case or script. For each, name the setting, two or more professional roles, the communication problem and one useful terminology focus. Include routine coordination, clarification, a complication, disagreement, handoff and follow-up where relevant. Choose scenarios that fit this profession, rather than forcing unsuitable situations into it.
2. Ask me to choose one scenario, or request all six in sequence. Stop and wait. Do not write all the scripts before I choose.
3. For the selected scenario, write an original fictional dialogue of 12-18 substantial speaking turns between two or three professionals. Name each role, establish an actual work problem, and let the exchange progress through questions, clarification, competing constraints and a credible next step or explicitly unresolved issue. Use the occupation's natural nomenclature and register. Avoid an interview between a teacher and a learner, generic small talk, and inserting a glossary definition into every reply. Explain specialist terms outside the dialogue instead.
4. After the script, explain five useful expressions in context, identify two grammar or register choices and ask three questions about the speakers' reasoning. Withhold the answers until I attempt them. Add one role-switch challenge.
5. Before presenting a script, check names, numbers, chronology, roles and terminology for consistency. Label invented facts and acknowledge uncertain specialist usage. If I requested all six, deliver one complete script at a time and wait for "next". Make each scenario materially different.

END PROMPT

Copy this prompt online or choose another lesson and support level