

ENGLISH FOR WORK / SEPTEMBER 2026

AI Development English



Teacher's guide

Teach practical workplace communication with specific cases, clear language targets, and a repeatable feedback process.

PRACTICE THAT TRANSFERS TO WORK

Read a realistic case. Study complete conversations. Find the right language, speak, get feedback, and try again.

Eight lessons | Intermediate to advanced | Independent or classroom study

For AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders.

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

All scenarios, figures, and model responses are fictional teaching material. Use invented details in practice.

[Open this course on English Ladder](#)

Your course at a glance

Use the lesson numbers to match this document with the website and the other guides.

01 Speaking the AI Development Stack

Explain a useful distinction in plain English.

02 LLMs, Transformers, Tokens, and Context

Clarify missing information before responding.

03 Data, Datasets, Labels, and Leakage

Separate observations, assumptions, and conclusions.

04 Retrieval, Embeddings, Vector Search, and RAG

Compare options using the same criteria and state the tradeoff.

05 Fine-Tuning, Alignment, and Adaptation

Clarify missing information before responding.

06 Evaluation, Benchmarks, and Regression

Distinguish completed work from pending work and explain its impact.

07 Inference, Latency, Cost, and Deployment

Negotiate scope or timing while making conditions explicit.

08 Safety, Security, Privacy, and Governance

Make a request with a clear action, purpose, and realistic timing.

Choose your pace

Quick practice: spend 15 minutes reading one conversation and rehearsing one scenario. Full lesson: allow 45-60 minutes for language work, three conversations, two speaking scenarios, and feedback.

Level support

B1 learners can use the frames and a partner. B2 learners can try a response before reading the model. C1 learners can add the harder follow-up and reduce preparation time. These are teaching suggestions, not certified CEFR level ratings.

After practicing: open the AI prompt workshop to choose a vocabulary, grammar, role-play, or new-dialogue activity.

A 60-minute lesson that works

Prepare the case and the learner pages. Keep the model response hidden until learners have attempted their own. Choose two field terms that are important to understanding the case.

0-5 minutes | Activate experience

Ask learners to name a comparable situation without sharing confidential details. Introduce the communication goal.

5-12 minutes | Read and sort facts

Read the case. Learners identify confirmed facts, unknowns, the audience, and the immediate communication need.

12-22 minutes | Notice and practice language

Explain the workshop pattern. Work through the vocabulary and editing example, then discuss both language checks and the reasons behind the choices.

22-35 minutes | 05 • Conversations

Read one ten-turn conversation aloud, notice a useful expression, then compare the other two. Identify each speaker's purpose and how the exchange ends.

35-47 minutes | 06 • Say it

Rehearse both speaking scenarios from the workbook. Give two minutes of preparation, switch roles, and use the success checks. Add the complication only after a successful first round.

47-55 minutes | Compare and revise

Reveal the model. Discuss alternative acceptable wording. Give one meaning correction and one language correction, then repeat.

55-60 minutes | Retrieve and transfer

Close the materials. Learners recall two expressions and identify one communication habit to use in another situation.

Adapt the session

For 45 minutes, study one complete conversation and rehearse one scenario; assign the others for oral follow-up. For 90 minutes, rehearse all three conversations and give both scenarios a second round. More time should produce more meaningful practice, not longer explanations.

LESSON 01 / FACILITATOR NOTES

Speaking the AI Development Stack

Explain a useful distinction in plain English.

CASE

A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Teach this move

Start with the reader's question. Explain one unfamiliar term with familiar words, then give a relevant example or contrast. Check understanding without asking only 'Do you understand?'.

Vocabulary to check: LLM, Foundation model, Inference, Throughput.

Model response

There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. The revised sentence names the action in familiar words. Specialist vocabulary is useful when it adds precision, but repeating abstract nouns can hide the meaning.

Second-round challenge

Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Language-check answers

1. C - The backlog is the work still waiting to be completed. This gives the meaning in familiar words.
2. B - How would you describe the next step in your own words? An open request lets the listener show their understanding.

05 • Conversations

1. Which component is failing? - An engineer and a product manager inspect an incorrect answer from a document assistant.
2. A queue is not a slow model - A serving engineer explains response timing to a new colleague.
3. One answer, two displays - A frontend developer and an AI engineer investigate a missing citation.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | The spinner stays on

The model service has returned an answer, but the interface still displays a loading indicator. A trace contains a successful service response; the frontend event log has not been checked.

Role A: You are the AI engineer. Explain the difference between completed inference and a completed user interaction.

Role B: You are the frontend developer. Ask which evidence is available and what you should inspect next.

Possible opening: The model returned an answer, but we haven't established why the interface is still waiting.

Success checks

- Distinguish the service response from the interface state.
- Do not claim the frontend cause is confirmed.
- Agree which log or event to inspect next.

Add a complication: The product manager asks you to describe the issue without using the words inference or payload.

LESSON 02 / FACILITATOR NOTES

LLMs, Transformers, Tokens, and Context

Clarify missing information before responding.

CASE

A prompt plus retrieved documents exceeds a model's context limit. A colleague proposes simply increasing the output length.

Teach this move

Name the exact word, fact, or requirement you need clarified. Ask one focused question at a time. Repeat your understanding and give the other person a chance to correct it.

Vocabulary to check: Token, Context window, Prompt, System prompt.

Model response

Do you mean a longer answer or more room for the input? Those are different limits. Let's check which messages and documents must fit in the context window.

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. After an introductory phrase such as 'Could you explain', the embedded question uses statement word order: "what this means".

Second-round challenge

The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Language-check answers

1. B - Could you confirm when the review starts? The embedded question uses subject then verb: "the review starts".
2. A - Has the report been approved, or is it ready for review? This question distinguishes two meanings of 'ready' that affect the next action.

05 • Conversations

1. More output space will not fit the input - Two engineers review a request rejected for exceeding the model's context limit.
2. The token count changed - A localization lead compares English and Japanese requests in the team's test environment.
3. Instructions and quoted text - A prompt designer and reviewer discuss a document summarization request.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | What can be shortened?

A test request is over the application's token budget. It contains a mandatory output schema, three repeated examples, and six retrieved passages. Relevance has not yet been assessed.

Role A: You maintain the prompt. Propose an order for reviewing redundant or irrelevant content.

Role B: You own the evaluation. Ask how the proposed shortening could affect correctness and follow-up questions.

Possible opening: Before removing constraints, let's check repeated examples and passage relevance.

Success checks

- Separate mandatory constraints from potentially redundant context.
- Ask for relevance evidence rather than guessing which passage is unnecessary.
- Propose a comparison using the same test questions.

Add a complication: A colleague wants to remove the output schema because it is the longest single instruction.

LESSON 03 / FACILITATOR NOTES

Data, Datasets, Labels, and Leakage

Separate observations, assumptions, and conclusions.

CASE

An evaluation scores 94%, but some evaluation examples also appear in the training data. A release meeting starts in an hour.

Teach this move

Say what the evidence shows, identify what is still unknown, and limit the conclusion to that evidence. Words such as 'may', 'suggests', and 'in this sample' are useful when the uncertainty is real. Do not soften an urgent, verified safety concern.

Vocabulary to check: Golden set, Eval, Regression, Pass rate.

Model response

The 94% result may be inflated because the datasets overlap. It does not yet show performance on unseen examples. I'll label this result and request a separate evaluation set.

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. The revision states the finding and its scope. It avoids turning limited evidence into a universal claim.

Second-round challenge

Someone asks for a yes-or-no answer when the evidence supports only a qualified answer. Be concise while preserving the uncertainty.

Language-check answers

1. A - The change may have contributed to the delay. 'May have contributed' presents a possibility without claiming the cause is established.
2. C - In this pilot, the observed result was ... This locates the finding in the setting actually studied.

05 · Conversations

1. The ninety-four percent result - An evaluation owner briefs a release lead before a meeting.
2. Duplicates with different filenames - A data engineer and evaluator examine two apparently separate dataset partitions.

3. Disagreement in the golden set - Two annotators and their lead discuss conflicting labels on a support-answer benchmark.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | A benchmark used as a tutorial

A developer pasted five evaluation questions into a team prompt tutorial. You do not yet know whether the tutorial examples were later included in fine-tuning data.

Role A: You own the benchmark. Explain the possible contamination and request provenance information.

Role B: You manage training data. Distinguish known reuse from unconfirmed inclusion in a training run.

Possible opening: The questions were reused in a tutorial, but we haven't traced whether they reached training.

Success checks

- State what reuse is confirmed.
- Avoid equating a missing filename match with evidence of independence.
- Agree how to trace content and qualify the current score.

Add a complication: The release lead asks you to call the benchmark clean because no matching filename was found.

LESSON 04 / FACILITATOR NOTES

Retrieval, Embeddings, Vector Search, and RAG

Compare options using the same criteria and state the tradeoff.

CASE

A retrieval experiment with reranking finds more relevant passages but adds 300 milliseconds to response time. No user study has been run.

Teach this move

Use the same time period, scope, and measurement basis for both options. Explain a difference with 'whereas' or 'while'. State what improves and what is given up; do not hide the weaker side of a recommendation.

Vocabulary to check: Embedding, Vector store, Reranker, RAG.

Model response

Reranking improved passage relevance in this test, but it added 300 milliseconds. We still need to check the effect on answer quality and the user experience before choosing a configuration.

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. 'Better' already has comparative meaning, so it does not need 'more'. Naming the criterion and tradeoff makes the comparison useful.

Second-round challenge

Your listener focuses only on the lowest price or fastest date. Explain one other relevant comparison criterion.

Language-check answers

1. C - Option A is faster, whereas option B costs less. 'Whereas' marks a contrast between two options.
2. B - An increase of one percentage point. Subtracting 6% from 7% gives one percentage point. The relative increase is about 16.7%.

05 • Conversations

1. Relevance versus response time - A retrieval engineer presents a reranking experiment to a product manager.
2. Similar is not necessarily relevant - A support specialist asks why vector search returned an unhelpful passage.

3. A citation that does not support the claim - An evaluator and engineer inspect a retrieval-augmented answer.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | A new index, a different comparison

Version A was tested on last month's document index. Version B uses a refreshed index and a reranker. Its relevance score is higher, but two things changed at once.

Role A: You are the experiment reviewer. Explain why the reranker's contribution is not isolated.

Role B: You implemented version B. Propose a controlled comparison without claiming the whole gain belongs to reranking.

Possible opening: The result improved, but the index refresh and reranker changed together.

Success checks

- Identify both changed components.
- Use cautious causal language.
- Agree on a comparison that keeps the index constant.

Add a complication: A stakeholder wants the slide title to say that reranking caused the full improvement.

LESSON 05 / FACILITATOR NOTES

Fine-Tuning, Alignment, and Adaptation

Clarify missing information before responding.

CASE

A summarizer uses the wrong format and sometimes cites an outdated refund policy. A colleague proposes fine-tuning as one solution to both problems.

Teach this move

Name the exact word, fact, or requirement you need clarified. Ask one focused question at a time. Repeat your understanding and give the other person a chance to correct it.

Vocabulary to check: Fine-tuning, SFT, LoRA, Adapter.

Model response

Could we separate the formatting problem from the outdated information? We can test a prompt change for the format and an updated retrieval source for policy facts before choosing an adaptation method.

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. After an introductory phrase such as 'Could you explain', the embedded question uses statement word order: "what this means".

Second-round challenge

The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Language-check answers

1. B - Could you confirm when the review starts? The embedded question uses subject then verb: "the review starts".
2. A - Has the report been approved, or is it ready for review? This question distinguishes two meanings of 'ready' that affect the next action.

05 • Conversations

1. One proposal, two failure types - A model engineer and policy-content owner review a summarizer's errors.
2. An adapter is still a change to evaluate - An engineer proposes a parameter-efficient adaptation experiment to a reviewer.

3. Good-looking examples, weak coverage - A training-data curator and engineer review candidate supervised examples.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | The format passes, the facts do not

An adapter's outputs pass JSON parsing and validation against the required schema on all twenty trial documents. Three extracted dates do not match their source documents. No held-out evaluation has been completed.

Role A: You developed the adapter. Report formatting success without implying factual correctness.

Role B: You approve evaluation plans. Ask for distinct checks of schema validity, source accuracy, and generalization.

Possible opening: The trial outputs all pass schema validation, but three dates are inaccurate and the held-out test is still pending.

Success checks

- Separate schema compliance from correct content.
- Mention the three inaccurate dates and incomplete evaluation.
- Agree on separate measures rather than one misleading success label.

Add a complication: Someone proposes calling the result one hundred percent accurate because every output passes schema validation.

LESSON 06 / FACILITATOR NOTES

Evaluation, Benchmarks, and Regression

Distinguish completed work from pending work and explain its impact.

CASE

A new model passes 88 of 100 test cases; the previous version passes 90. The new version performs better on long documents but worse on short requests.

Teach this move

Lead with the current state, then the important change or open item. Use completed-action language only for completed work. Add a next update point when it is known; do not invent a date to make the message sound complete.

Vocabulary to check: Eval, Benchmark, Pass rate, Regression.

Model response

The overall pass rate fell from 90% to 88%. Long-document results improved, while short-request results declined. I'll show both groups so the release decision reflects the tradeoff.

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. The revision gives observable progress. A positive label alone does not tell the reader what is complete or still needed.

Second-round challenge

The listener asks you to reduce the update to two sentences. Preserve the most important completed and pending items.

Language-check answers

1. A - The team has completed the review. 'Has completed' states that the action is finished.
2. C - Delivery is planned for Friday; the carrier has not confirmed it. This explicitly labels the date as planned and identifies the missing confirmation.

05 · Conversations

1. Two fewer passes, one stronger slice - An evaluator briefs a release manager on a model comparison.
2. Did the model change or the grader? - Two evaluation engineers investigate an unexpected jump in pass rate.

3. A serious failure hidden in the average - A quality reviewer discusses two failure categories with an evaluation owner.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | The denominator changed

A dashboard shows eighty passes out of ninety evaluated cases. Ten of the original hundred cases failed to run. A colleague compares the displayed percentage with last week's ninety passes out of one hundred.

Role A: You own the dashboard. Explain the missing cases and why the displayed rates are not directly comparable.

Role B: You are the release lead. Ask what happened to the ten cases and request a like-for-like comparison.

Possible opening: This week's rate excludes ten cases that did not run, so the denominator has changed.

Success checks

- State both denominators accurately.
- Distinguish an unexecuted case from a passing case.
- Agree to rerun or explicitly report the missing results.

Add a complication: A stakeholder asks you to count the missing cases as passes because no incorrect answer was recorded.

LESSON 07 / FACILITATOR NOTES

Inference, Latency, Cost, and Deployment

Negotiate scope or timing while making conditions explicit.

CASE

A team can reduce average response time from four seconds to two by using a smaller model. Complex-answer quality has not yet been checked.

Teach this move

Acknowledge the goal and explain the available choices. Connect an offer to its actual conditions. A useful response makes room for discussion without promising approval or resources you do not control.

Vocabulary to check: Inference, Latency, Batching, Fallback.

Model response

The smaller model is faster in this test. If it meets our quality criteria on complex requests, we could use it more widely. Until then, I'd keep the rollout limited.

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. For an ordinary future condition, the "if" clause uses present tense. The other clause can express the future consequence or possibility.

Second-round challenge

The preferred option is unavailable. Offer an alternative and clearly state any approval still needed.

Language-check answers

1. C - If the scope changes, we will review the schedule. The "if" clause uses present tense; the main clause states the future response.
2. B - We can assess the addition and explain its effect on the agreed scope. This offers a useful next step while keeping the approval and scope question open.

05 • Conversations

1. A deployment offer with conditions - An infrastructure engineer and product lead negotiate the scope of a smaller-model trial. The release owner must approve any deployment.
2. Cheap requests, expensive retries - A finance partner asks an inference engineer about a lower-cost serving configuration.
3. The fallback needs its own failure test - A deployment engineer reviews a fallback plan with the on-call lead.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | Streaming looks faster

A prototype displays its first text sooner. Total answer completion time has not been measured. The current slide says inference is twice as fast, although only time to first displayed text was measured.

Role A: You are the serving engineer. Offer to publish the narrower first-text result now, or defer the broader comparison until total completion time is measured.

Role B: You are the product manager. Negotiate which claim to publish and what evidence is required before expanding it.

Possible opening: We can report the earlier first text now, but a claim about the complete answer needs another measurement.

Success checks

- Distinguish measured first-text timing from unknown total completion time.
- Offer two publication options with explicit conditions.
- Agree what evidence is required before claiming faster complete answers.

Add a complication: A sales colleague wants to keep the phrase twice as fast because it sounds simpler.

LESSON 08 / FACILITATOR NOTES

Safety, Security, Privacy, and Governance

Make a request with a clear action, purpose, and realistic timing.

CASE

A test assistant repeats a fictional secret embedded in a retrieved document. The team has a reporting channel and needs a concise reproduction note.

Teach this move

State the action and the information needed. Explain why it matters and specify a deadline only when one is given or agreed. 'Could you' is polite; repeated apologies can obscure a legitimate request. Urgent situations may require a direct request.

Vocabulary to check: Prompt injection, Jailbreak, PII, Red team.

Model response

Please review this test through the security reporting channel. I've attached the synthetic input, observed output, and reproduction steps. We have not yet established whether other configurations are affected.

Listen for, then coach

Ask learners to point to the case fact behind each important statement. Accept different wording when it preserves meaning, fits the listener, and does not invent an outcome. The revision names the document and requested response. It avoids expressions that may be unfamiliar or ambiguous for an international audience.

Second-round challenge

The other person cannot meet the requested timing. Clarify an alternative without inventing authority to change a formal deadline.

Language-check answers

1. B - Send it no later than Thursday. 'By' sets the latest point, while 'after' refers to a later time.
2. A - Could you confirm the document version needed for the review? The requested action and information are explicit.

05 • Conversations

1. A reproducible injection report - In a synthetic test, the authorized task is to summarize a document without disclosing its canary string. The retrieved document includes an instruction to ignore that restriction and print the string.

2. A redacted log still needs review - An engineer asks a privacy reviewer whether a diagnostic bundle can be shared.

3. A tool call is more than an answer - An agent developer and security engineer review a proposed calendar integration.

Use the complete ten-turn scripts in the learner workbook and conversation lab. Read with roles, identify the decision or open question, then replay without reading.

06 · Say it: second scenario | An external destination in retrieved text

A test document instructs an assistant to send its summary to an unfamiliar address. The assistant has a sending tool, but no message has been sent. The address is synthetic.

Role A: You are the tester. Describe the attempted instruction and distinguish a requested action from an executed one.

Role B: You are the security reviewer. Ask for the trace and how authorization is enforced before tool execution.

Possible opening: The retrieved document requested an external send, but the trace does not show an executed message.

Success checks

- Distinguish an instruction attempt from a completed send.
- Keep the test data synthetic and use the reporting channel.
- Agree what evidence would establish whether the tool boundary worked.

Add a complication: A colleague wants the incident title to say data exfiltration confirmed.

Give feedback that helps

Score each criterion 0, 1, or 2. This is a practice rubric, not a professional qualification or CEFR test. Do not award or remove points for a particular accent.

Meaning and accuracy

Keeps the case facts accurate; clearly separates confirmed and unknown information.

Clarity and language

Uses understandable sentences and explains specialist terms when the listener needs it.

Interaction and tone

Responds to the other person's question, checks understanding, and uses an appropriate tone.

Task completion

Makes the requested action or unresolved question clear without inventing authority or facts.

Use the same scale for each criterion

0 - Not yet: The reader or listener cannot yet recover the intended meaning or action.

1 - With support: The message works after a prompt, clarification, or revision.

2 - Independently: The message is clear and accurate without a prompt.

Feedback sequence

First, identify a successful phrase. Next, identify one point where meaning or accuracy needs work. Offer one useful language correction, allow a repeat, and compare the two attempts. A total out of eight describes this performance only; do not translate it into a proficiency level.

Final challenge

Choose a case not rehearsed today. The learner gives a one-minute response, answers two follow-up questions, and writes a 70-110 word message. Add the lesson's second-round challenge. Compare the result with an earlier attempt using the four criteria.

Use the language responsibly

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Meaning before performance

Use specialist terms only when they help the listener. Explain unfamiliar words, preserve uncertainty, and ask when an instruction or responsibility is unclear. The model responses illustrate language; they do not authorize professional actions.

Sources and further reading

The cases, workshops, and explanations are original teaching material. These references provide language frameworks and selected professional context. References were checked in September 2026; consult current local guidance for actual work.

Digital.gov: plain-language guidance

<https://digital.gov/guides/plain-language>

Council of Europe: CEFR descriptors

<https://www.coe.int/en/web/common-european-framework-reference-languages/cefr-descriptors>

Your next step

Choose one expression you can use in a genuine work situation. Rehearse it with fictional details, then adapt the wording to your role and audience.