

ENGLISH FOR WORK / SEPTEMBER 2026

AI Development English



Learner workbook

Work with realistic cases, develop your own response, and use the answer section after you have tried the tasks.

PRACTICE THAT TRANSFERS TO WORK

Read a realistic case. Study complete conversations. Find the right language, speak, get feedback, and try again.

Eight lessons | Intermediate to advanced | Independent or classroom study

For AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders.

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

All scenarios, figures, and model responses are fictional teaching material. Use invented details in practice.

[Open this course on English Ladder](#)

Your course at a glance

Use the lesson numbers to match this document with the website and the other guides.

01 Speaking the AI Development Stack

Explain a useful distinction in plain English.

02 LLMs, Transformers, Tokens, and Context

Clarify missing information before responding.

03 Data, Datasets, Labels, and Leakage

Separate observations, assumptions, and conclusions.

04 Retrieval, Embeddings, Vector Search, and RAG

Compare options using the same criteria and state the tradeoff.

05 Fine-Tuning, Alignment, and Adaptation

Clarify missing information before responding.

06 Evaluation, Benchmarks, and Regression

Distinguish completed work from pending work and explain its impact.

07 Inference, Latency, Cost, and Deployment

Negotiate scope or timing while making conditions explicit.

08 Safety, Security, Privacy, and Governance

Make a request with a clear action, purpose, and realistic timing.

Choose your pace

Quick practice: spend 15 minutes reading one conversation and rehearsing one scenario. Full lesson: allow 45-60 minutes for language work, three conversations, two speaking scenarios, and feedback.

Level support

B1 learners can use the frames and a partner. B2 learners can try a response before reading the model. C1 learners can add the harder follow-up and reduce preparation time. These are teaching suggestions, not certified CEFR level ratings.

After practicing: open the AI prompt workshop, or use the prompts at the end of this guide.

LESSON 01 / READ AND NOTICE

Speaking the AI Development Stack

Explain a useful distinction in plain English.

SITUATION

A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

LLM - Large language model; a model trained to process and generate language-like sequences.

Foundation model - A broadly trained model adapted to many downstream tasks.

Inference - Running a trained model to produce an output.

Throughput - How many requests a system can handle in a period of time.

3. Notice the language

Start with the reader's question. Explain one unfamiliar term with familiar words, then give a relevant example or contrast. Check understanding without asking only 'Do you understand?'.

In this context, ... means

The difference is that

How would you explain that distinction to a colleague?

Compare this wording: We need to operationalize the implementation.

LESSON 01 / PRACTICE AND PRODUCE

Make a complex point clear

4. Check the language

1. Which sentence explains 'backlog' rather than repeating it?
 - A. The backlog is the work completed during the last period.
 - B. The backlog is the maximum work the team can finish in a period.
 - C. The backlog is the work still waiting to be completed.
2. Which follow-up best checks whether your explanation was clear?
 - A. Was the explanation detailed enough for you?
 - B. How would you describe the next step in your own words?
 - C. Would you like me to repeat the explanation?

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 01 / 05 • CONVERSATIONS / 1 OF 3

Which component is failing?

An engineer and a product manager inspect an incorrect answer from a document assistant.

01 • Product manager: The AI gave the wrong answer. Should we replace the model?

02 • Engineer: First, let's separate retrieval, answer generation, and what the interface displays.

03 • Product manager: The displayed answer refers to last year's handbook. Which layer controls that?

04 • Engineer: Retrieval supplies passages. The language model uses them to draft an answer.

05 • Product manager: So an old passage could produce an answer that sounds current?

06 • Engineer: Yes. We need the retrieved passages before attributing this example to the model.

07 • Product manager: The trace shows last year's handbook among the returned documents.

08 • Engineer: Then let's inspect document versioning and retrieval before changing the foundation model.

09 • Product manager: I'll attach the question, displayed answer, and document version to the issue.

10 • Engineer: Good. I'll trace that request and distinguish confirmed findings from possible causes.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 01 / 05 • CONVERSATIONS / 2 OF 3

A queue is not a slow model

A serving engineer explains response timing to a new colleague.

01 • Colleague: Users waited twelve seconds, but the inference log records only four.

02 • Serving engineer: The request also spent eight seconds waiting in the serving queue.

03 • Colleague: Does inference mean the entire period from clicking Send to seeing an answer?

04 • Serving engineer: Here, inference means running the model. End-to-end latency includes the other stages.

05 • Colleague: Then our dashboard label is making two different measures look interchangeable.

06 • Serving engineer: Exactly. Let's label queue time, model execution, and total response time separately.

07 • Colleague: Would increasing throughput mean every individual request becomes faster?

08 • Serving engineer: Not necessarily. Throughput measures completed work per period, not one request's duration.

09 • Colleague: I'll revise the labels and include both measures in the incident summary.

10 • Serving engineer: Include the measurement boundaries too, so the next comparison uses the same definition.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 01 / 05 • CONVERSATIONS / 3 OF 3

One answer, two displays

A frontend developer and an AI engineer investigate a missing citation.

01 • Frontend developer: The answer on screen has no citation, so retrieval must have failed.

02 • AI engineer: Let's check the response payload before drawing that conclusion.

03 • Frontend developer: The payload contains a passage identifier and a citation URL.

04 • AI engineer: Then the backend returned citation information. The display may be dropping it.

05 • Frontend developer: The new interface expects a citations array, but this response uses source.

06 • AI engineer: That looks like a contract mismatch between the service and the interface.

07 • Frontend developer: I'll reproduce it with a saved response rather than another live model call.

08 • AI engineer: Good. That isolates rendering from retrieval and generation variability.

09 • Frontend developer: Once the mapping is fixed, should we also test empty citation lists?

10 • AI engineer: Yes. Check both cases and avoid treating missing evidence as a display-only problem.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 01 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Speaking the AI Development Stack

A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Second round: Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Model for scenario 1: There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

Scenario 2 | The spinner stays on

The model service has returned an answer, but the interface still displays a loading indicator. A trace contains a successful service response; the frontend event log has not been checked.

Role A: You are the AI engineer. Explain the difference between completed inference and a completed user interaction.

Role B: You are the frontend developer. Ask which evidence is available and what you should inspect next.

Possible opening: The model returned an answer, but we haven't established why the interface is still waiting.

Success checks

- Distinguish the service response from the interface state.
- Do not claim the frontend cause is confirmed.
- Agree which log or event to inspect next.

Add a complication: The product manager asks you to describe the issue without using the words inference or payload.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

LESSON 02 / READ AND NOTICE

LLMs, Transformers, Tokens, and Context

Clarify missing information before responding.

SITUATION

A prompt plus retrieved documents exceeds a model's context limit. A colleague proposes simply increasing the output length.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

Token - A unit of text processed by the model; token count affects cost, context, and latency.

Context window - The amount of input and generated text the model can consider in one request.

Prompt - The instructions, examples, user request, and context given to a model.

System prompt - High-priority instructions that guide model behavior inside an application.

3. Notice the language

Name the exact word, fact, or requirement you need clarified. Ask one focused question at a time. Repeat your understanding and give the other person a chance to correct it.

When you say ..., do you mean ...?

Could you clarify which ...?

So, my understanding is Is that right?

Compare this wording: Could you explain what does this mean?

LESSON 02 / PRACTICE AND PRODUCE

Ask a precise question

4. Check the language

1. Which indirect question has the usual word order?
 - A. Could you confirm when starts the review?
 - B. Could you confirm when the review starts?
 - C. Could you confirm when does the review start?
2. Someone says the report is 'ready.' Which question best clarifies its approval status?
 - A. Has the report been approved, or is it ready for review?
 - B. Can you send the report to me this afternoon?
 - C. Who will present the report at the meeting?

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 02 / 05 • CONVERSATIONS / 1 OF 3

More output space will not fit the input

Two engineers review a request rejected for exceeding the model's context limit.

01 • Developer: The request is too long. Could we just increase the output allowance?

02 • AI engineer: That won't shorten the prompt or the retrieved documents already in the request.

03 • Developer: I thought the context window meant the answer's maximum length.

04 • AI engineer: It is the model's available context, not simply the number of answer tokens.

05 • Developer: Our application reserves answer space before adding retrieved passages.

06 • AI engineer: Then we should inspect that budget and count the assembled request.

07 • Developer: Could I delete the oldest conversation messages until it fits?

08 • AI engineer: Only after checking whether they contain constraints the current question still needs.

09 • Developer: I'll measure each part and propose a documented truncation rule.

10 • AI engineer: Test that rule on follow-up questions, not only standalone prompts.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 02 / 05 • CONVERSATIONS / 2 OF 3

The token count changed

A localization lead compares English and Japanese requests in the team's test environment.

01 • Localization lead: Both requests have about the same character count. Why do their token counts differ?

02 • Engineer: The tokenizer doesn't assign exactly one token to each character or word.

03 • Localization lead: Can I use the English average to estimate the Japanese requests?

04 • Engineer: We should measure those requests with the tokenizer used by our deployed model.

05 • Localization lead: The Japanese sample includes product codes and mixed punctuation as well.

06 • Engineer: Include those in the test set. They are part of our actual input.

07 • Localization lead: Would a word limit in the interface still protect the context budget?

08 • Engineer: It would be an approximation, not a guarantee that the assembled request fits.

09 • Localization lead: I'll collect representative samples and flag the largest assembled requests.

10 • Engineer: I'll report token counts by sample and keep the tokenizer version in the results.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 02 / 05 • CONVERSATIONS / 3 OF 3

Instructions and quoted text

A prompt designer and reviewer discuss a document summarization request.

01 • Reviewer: This quoted document tells the assistant to ignore the requested output format.

02 • Prompt designer: That sentence is document content, not an instruction we intend to authorize.

03 • Reviewer: Is placing our format rule in the system prompt sufficient proof of protection?

04 • Prompt designer: No. We still need to test whether the assistant follows the intended boundary.

05 • Reviewer: Let's include documents that contain instruction-like language in the evaluation.

06 • Prompt designer: Agreed. I'll keep document text clearly delimited in the request structure.

07 • Reviewer: Should the summary mention the instruction-like sentence if it is relevant?

08 • Prompt designer: It may describe the sentence as content without treating it as a command.

09 • Reviewer: I'll check both the output format and whether any unauthorized action occurred.

10 • Prompt designer: Record the full test conditions so a passing example doesn't become a blanket claim.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 02 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | LLMs, Transformers, Tokens, and Context

A prompt plus retrieved documents exceeds a model's context limit. A colleague proposes simply increasing the output length.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Second round: The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Model for scenario 1: Do you mean a longer answer or more room for the input? Those are different limits. Let's check which messages and documents must fit in the context window.

Scenario 2 | What can be shortened?

A test request is over the application's token budget. It contains a mandatory output schema, three repeated examples, and six retrieved passages. Relevance has not yet been assessed.

Role A: You maintain the prompt. Propose an order for reviewing redundant or irrelevant content.

Role B: You own the evaluation. Ask how the proposed shortening could affect correctness and follow-up questions.

Possible opening: Before removing constraints, let's check repeated examples and passage relevance.

Success checks

- Separate mandatory constraints from potentially redundant context.
- Ask for relevance evidence rather than guessing which passage is unnecessary.
- Propose a comparison using the same test questions.

Add a complication: A colleague wants to remove the output schema because it is the longest single instruction.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

LESSON 03 / READ AND NOTICE

Data, Datasets, Labels, and Leakage

Separate observations, assumptions, and conclusions.

SITUATION

An evaluation scores 94%, but some evaluation examples also appear in the training data. A release meeting starts in an hour.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

Golden set - Curated examples used repeatedly to test important behavior.

Eval - A test or evaluation suite for model or system behavior.

Regression - A behavior that gets worse after a change.

Pass rate - The percentage of eval cases meeting the success criterion.

3. Notice the language

Say what the evidence shows, identify what is still unknown, and limit the conclusion to that evidence. Words such as 'may', 'suggests', and 'in this sample' are useful when the uncertainty is real. Do not soften an urgent, verified safety concern.

The available evidence shows

This may indicate ..., but

We have not yet established whether

Compare this wording: The pilot proves this will work everywhere.

LESSON 03 / PRACTICE AND PRODUCE

Match certainty to evidence

4. Check the language

1. Which sentence preserves uncertainty about an unconfirmed cause?
 - A. The change may have contributed to the delay.
 - B. The change must have caused the delay.
 - C. The change definitely caused the delay.
2. A result comes from a small pilot in one location. Which phrase accurately limits the claim?
 - A. In similar locations, we can now expect the same result ...
 - B. Across the service, the observed result was ...
 - C. In this pilot, the observed result was ...

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 03 / 05 • CONVERSATIONS / 1 OF 3

The ninety-four percent result

An evaluation owner briefs a release lead before a meeting.

01 • Release lead: Can I present the ninety-four percent score as evidence that we're ready?

02 • Evaluation owner: Not without a caveat. Some evaluation examples also occur in training data.

03 • Release lead: Does that mean every answer in the evaluation is invalid?

04 • Evaluation owner: No, but the overlap compromises the independence of the reported evaluation.

05 • Release lead: Can you calculate a clean score before the meeting in an hour?

06 • Evaluation owner: I can identify known overlaps. I can't promise a completed independent rerun yet.

07 • Release lead: Then what should the slide say about the current pass rate?

08 • Evaluation owner: Label it provisional and state that training-evaluation overlap was detected.

09 • Release lead: I'll separate the current result from any claim about unseen examples.

10 • Evaluation owner: I'll bring the overlap evidence and the proposed replacement evaluation plan.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 03 / 05 • CONVERSATIONS / 2 OF 3

Duplicates with different filenames

A data engineer and evaluator examine two apparently separate dataset partitions.

01 • Evaluator: The file identifiers differ, so I assumed these examples were independent.

02 • Data engineer: The underlying text can still be duplicated even when identifiers change.

03 • Evaluator: This question has different punctuation but the same reference answer.

04 • Data engineer: Let's check exact duplicates and near-duplicates, and record how we identify each.

05 • Evaluator: Should we remove every example that discusses the same subject?

06 • Data engineer: Shared subject matter isn't automatically duplication. We need a defensible matching rule.

07 • Evaluator: I'll review borderline matches rather than silently deleting them.

08 • Data engineer: Keep the original split and the filtered version so the change is auditable.

09 • Evaluator: The report will show the rule, removed counts, and remaining evaluation size.

10 • Data engineer: Good. Don't compare the new percentage with the old one without explaining the changed set.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 03 / 05 • CONVERSATIONS / 3 OF 3

Disagreement in the golden set

Two annotators and their lead discuss conflicting labels on a support-answer benchmark.

01 • Annotator: We disagree on this answer. I marked it correct because the main fact is right.

02 • Annotation lead: What did the other labeler identify as the error?

03 • Annotator: The answer omits a required condition from the reference policy.

04 • Annotation lead: Does our rubric say whether that omission changes the correctness label?

05 • Annotator: It says complete, but it doesn't define which details are essential.

06 • Annotation lead: Then resolve the rubric ambiguity before treating either label as authoritative.

07 • Annotator: Should we send the model's response back through training immediately?

08 • Annotation lead: No. First record the adjudication and preserve the evaluation boundary.

09 • Annotator: I'll draft a specific rule and bring comparable examples for review.

10 • Annotation lead: Then recheck affected labels consistently and version the golden set.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 03 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Data, Datasets, Labels, and Leakage

An evaluation scores 94%, but some evaluation examples also appear in the training data. A release meeting starts in an hour.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to tell a colleague what is known. Ask which part is confirmed and which part is an assumption. Challenge one statement that sounds more certain than the case supports.

Second round: Someone asks for a yes-or-no answer when the evidence supports only a qualified answer. Be concise while preserving the uncertainty.

Model for scenario 1: The 94% result may be inflated because the datasets overlap. It does not yet show performance on unseen examples. I'll label this result and request a separate evaluation set.

Scenario 2 | A benchmark used as a tutorial

A developer pasted five evaluation questions into a team prompt tutorial. You do not yet know whether the tutorial examples were later included in fine-tuning data.

Role A: You own the benchmark. Explain the possible contamination and request provenance information.

Role B: You manage training data. Distinguish known reuse from unconfirmed inclusion in a training run.

Possible opening: The questions were reused in a tutorial, but we haven't traced whether they reached training.

Success checks

- State what reuse is confirmed.
- Avoid equating a missing filename match with evidence of independence.
- Agree how to trace content and qualify the current score.

Add a complication: The release lead asks you to call the benchmark clean because no matching filename was found.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

LESSON 04 / READ AND NOTICE

Retrieval, Embeddings, Vector Search, and RAG

Compare options using the same criteria and state the tradeoff.

SITUATION

A retrieval experiment with reranking finds more relevant passages but adds 300 milliseconds to response time. No user study has been run.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

Embedding - A vector representation used for similarity search, clustering, classification, and related tasks.

Vector store - A database or index for storing and searching embeddings.

Reranker - A model or step that reorders retrieved results for relevance.

RAG - Retrieval-augmented generation: retrieve relevant context, then generate an answer using it.

3. Notice the language

Use the same time period, scope, and measurement basis for both options. Explain a difference with 'whereas' or 'while'. State what improves and what is given up; do not hide the weaker side of a recommendation.

Both options ..., but

Option A ..., whereas option B

The comparison depends on

Compare this wording: This option is more better.

LESSON 04 / PRACTICE AND PRODUCE

Explain a fair comparison

4. Check the language

1. Which sentence expresses a clear contrast?
 - A. Option A is faster because option B costs less.
 - B. Option A is faster, so option B costs less.
 - C. Option A is faster, whereas option B costs less.
2. A rate increases from 6% to 7%. What is the absolute difference?
 - A. An increase of seven percentage points.
 - B. An increase of one percentage point.
 - C. An increase of one percent relative to the original rate.

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 04 / 05 • CONVERSATIONS / 1 OF 3

Relevance versus response time

A retrieval engineer presents a reranking experiment to a product manager.

01 • Product manager: The reranker finds better passages. Can we announce a better user experience?

02 • Retrieval engineer: We measured passage relevance, not user experience. It also adds three hundred milliseconds.

03 • Product manager: Is that enough delay for users to notice?

04 • Retrieval engineer: We haven't run a user study, so I wouldn't claim either way.

05 • Product manager: Then how should we describe the benefit in the experiment summary?

06 • Retrieval engineer: Higher measured retrieval relevance, with additional response time in this test.

07 • Product manager: Could we compare answers with and without reranking on the same questions?

08 • Retrieval engineer: Yes. That would help establish whether the better passages improve answer quality.

09 • Product manager: Let's include the latency distribution as well as the average.

10 • Retrieval engineer: I'll keep the query set fixed and document the comparison conditions.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 04 / 05 • CONVERSATIONS / 2 OF 3

Similar is not necessarily relevant

A support specialist asks why vector search returned an unhelpful passage.

01 • Support specialist: This passage mentions refunds repeatedly, but it doesn't answer the deadline question.

02 • Search engineer: Semantic similarity alone doesn't guarantee that a passage contains the required detail.

03 • Support specialist: Would asking for more results solve the problem?

04 • Search engineer: It might recover a useful passage, but it can also add irrelevant context.

05 • Support specialist: The policy date is important here. Can we filter by version?

06 • Search engineer: We can test that if the index has reliable version metadata.

07 • Support specialist: Some older documents don't have that field populated.

08 • Search engineer: Then missing metadata needs explicit handling; a filter isn't a complete fix.

09 • Support specialist: I'll supply examples where the current policy is missed or displaced.

10 • Search engineer: I'll compare retrieval results and identify whether indexing, filtering, or ranking is responsible.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 04 / 05 • CONVERSATIONS / 3 OF 3

A citation that does not support the claim

An evaluator and engineer inspect a retrieval-augmented answer.

01 • Evaluator: The answer includes a citation, but the cited paragraph doesn't support its conclusion.

02 • Engineer: Let's separate citation presence from whether the claim is grounded in that passage.

03 • Evaluator: The passage says requests are reviewed; the answer says approval is automatic.

04 • Engineer: That changes the meaning. A valid URL doesn't make the statement supported.

05 • Evaluator: Should I score this as a retrieval failure or a generation failure?

06 • Engineer: Check whether a supporting passage was available in the retrieved context first.

07 • Evaluator: The correct paragraph was present, and it explicitly described a review.

08 • Engineer: Record that evidence. This example points to how the answer used the context.

09 • Evaluator: I'll preserve the claim, cited span, and retrieved context in the failure report.

10 • Engineer: Then we can test a remedy against the actual failure, not just citation formatting.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 04 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Retrieval, Embeddings, Vector Search, and RAG

A retrieval experiment with reranking finds more relevant passages but adds 300 milliseconds to response time. No user study has been run.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You favor one option but have not checked whether the scope is comparable. Ask about one difference, then explain the tradeoff you are willing to accept.

Second round: Your listener focuses only on the lowest price or fastest date. Explain one other relevant comparison criterion.

Model for scenario 1: Reranking improved passage relevance in this test, but it added 300 milliseconds. We still need to check the effect on answer quality and the user experience before choosing a configuration.

Scenario 2 | A new index, a different comparison

Version A was tested on last month's document index. Version B uses a refreshed index and a reranker. Its relevance score is higher, but two things changed at once.

Role A: You are the experiment reviewer. Explain why the reranker's contribution is not isolated.

Role B: You implemented version B. Propose a controlled comparison without claiming the whole gain belongs to reranking.

Possible opening: The result improved, but the index refresh and reranker changed together.

Success checks

- Identify both changed components.
- Use cautious causal language.
- Agree on a comparison that keeps the index constant.

Add a complication: A stakeholder wants the slide title to say that reranking caused the full improvement.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

LESSON 05 / READ AND NOTICE

Fine-Tuning, Alignment, and Adaptation

Clarify missing information before responding.

SITUATION

A summarizer uses the wrong format and sometimes cites an outdated refund policy. A colleague proposes fine-tuning as one solution to both problems.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

Fine-tuning - Updating model weights on task- or domain-specific data.

SFT - Supervised fine-tuning with input-output examples.

LoRA - Low-rank adaptation; a parameter-efficient fine-tuning method.

Adapter - A small trainable module inserted into or attached to a pretrained model.

3. Notice the language

Name the exact word, fact, or requirement you need clarified. Ask one focused question at a time. Repeat your understanding and give the other person a chance to correct it.

When you say ..., do you mean ...?

Could you clarify which ...?

So, my understanding is Is that right?

Compare this wording: Could you explain what does this mean?

LESSON 05 / PRACTICE AND PRODUCE

Ask a precise question

4. Check the language

1. Which indirect question has the usual word order?
 - A. Could you confirm when starts the review?
 - B. Could you confirm when the review starts?
 - C. Could you confirm when does the review start?
2. Someone says the report is 'ready.' Which question best clarifies its approval status?
 - A. Has the report been approved, or is it ready for review?
 - B. Can you send the report to me this afternoon?
 - C. Who will present the report at the meeting?

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 05 / 05 • CONVERSATIONS / 1 OF 3

One proposal, two failure types

A model engineer and policy-content owner review a summarizer's errors.

01 • Model engineer: We could fine-tune the summarizer to fix the formatting and policy errors.

02 • Content owner: The formatting is inconsistent, but the cited refund policy is also outdated.

03 • Model engineer: You're distinguishing output behavior from the freshness of the source information?

04 • Content owner: Exactly. Would training on examples ensure that later policy changes reach the assistant?

05 • Model engineer: No. We still need a reliable path for maintaining current source content.

06 • Content owner: Then let's diagnose the two failures separately before choosing an intervention.

07 • Model engineer: I'll test formatting constraints and examples against a held-out set.

08 • Content owner: I'll check which policy versions are indexed and how replacements are published.

09 • Model engineer: If we try supervised fine-tuning, we'll measure its effect on formatting explicitly.

10 • Content owner: And we won't report that as proof that policy freshness is solved.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 05 / 05 • CONVERSATIONS / 2 OF 3

An adapter is still a change to evaluate

An engineer proposes a parameter-efficient adaptation experiment to a reviewer.

01 • Engineer: A LoRA adapter would let us test adaptation without updating every model parameter.

02 • Reviewer: Does that mean we can skip the full regression evaluation?

03 • Engineer: No. A smaller training update can still change important model behavior.

04 • Reviewer: What outcome are you targeting in this experiment?

05 • Engineer: Consistent extraction into our approved schema, not broader improvements in reasoning.

06 • Reviewer: Then include schema validity and whether extracted values preserve the source meaning.

07 • Engineer: I'll also keep examples outside the training set for evaluation.

08 • Reviewer: How will you tell whether the adapter improves on the current prompt-only version?

09 • Engineer: I'll use the same held-out inputs and report both results side by side.

10 • Reviewer: Good. Keep the objective narrow and record any regressions outside the target behavior.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 05 / 05 • CONVERSATIONS / 3 OF 3

Good-looking examples, weak coverage

A training-data curator and engineer review candidate supervised examples.

01 • Curator: These examples all produce neat summaries, so the dataset looks ready.

02 • Engineer: Do they include ambiguous inputs, missing fields, and conflicting source passages?

03 • Curator: Mostly straightforward documents. We removed the difficult ones during cleanup.

04 • Engineer: Then the examples may not represent the situations the deployed summarizer faces.

05 • Curator: Should we restore every rejected document?

06 • Engineer: Not automatically. Check quality and label the intended behavior for each difficult case.

07 • Curator: Some reviewers filled in missing facts to make the target summary complete.

08 • Engineer: We should correct that. The target should not teach unsupported completion.

09 • Curator: I'll separate genuine omissions from cases where the source itself lacks information.

10 • Engineer: Then document the coverage gaps before deciding whether this set is suitable for training.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 05 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Fine-Tuning, Alignment, and Adaptation

A summarizer uses the wrong format and sometimes cites an outdated refund policy. A colleague proposes fine-tuning as one solution to both problems.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to know exactly what information is missing. Give one answer only after your partner asks a focused question; then ask them to confirm their understanding.

Second round: The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

Model for scenario 1: Could we separate the formatting problem from the outdated information? We can test a prompt change for the format and an updated retrieval source for policy facts before choosing an adaptation method.

Scenario 2 | The format passes, the facts do not

An adapter's outputs pass JSON parsing and validation against the required schema on all twenty trial documents. Three extracted dates do not match their source documents. No held-out evaluation has been completed.

Role A: You developed the adapter. Report formatting success without implying factual correctness.

Role B: You approve evaluation plans. Ask for distinct checks of schema validity, source accuracy, and generalization.

Possible opening: The trial outputs all pass schema validation, but three dates are inaccurate and the held-out test is still pending.

Success checks

- Separate schema compliance from correct content.
- Mention the three inaccurate dates and incomplete evaluation.
- Agree on separate measures rather than one misleading success label.

Add a complication: Someone proposes calling the result one hundred percent accurate because every output passes schema validation.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

LESSON 06 / READ AND NOTICE

Evaluation, Benchmarks, and Regression

Distinguish completed work from pending work and explain its impact.

SITUATION

A new model passes 88 of 100 test cases; the previous version passes 90. The new version performs better on long documents but worse on short requests.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

Eval - A test or evaluation suite for model or system behavior.

Benchmark - A standardized test used to compare systems, often imperfect for a product use case.

Pass rate - The percentage of eval cases meeting the success criterion.

Regression - A behavior that gets worse after a change.

3. Notice the language

Lead with the current state, then the important change or open item. Use completed-action language only for completed work. Add a next update point when it is known; do not invent a date to make the message sound complete.

We have completed ...; ... remains open.

As of ..., the status is

The next update will cover

Compare this wording: Everything is progressing well.

LESSON 06 / PRACTICE AND PRODUCE

Give a useful status update

4. Check the language

1. Which sentence says a review is finished?
 - A. The team has completed the review.
 - B. The team is completing the review.
 - C. The team plans to complete the review.
2. Which update distinguishes a plan from a confirmed fact?
 - A. Delivery is confirmed for Friday because we planned it.
 - B. The delivery plan proves that arrival will be Friday.
 - C. Delivery is planned for Friday; the carrier has not confirmed it.

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 06 / 05 • CONVERSATIONS / 1 OF 3

Two fewer passes, one stronger slice

An evaluator briefs a release manager on a model comparison.

01 • Release manager: The new model passes eighty-eight cases; the old one passes ninety. Reject it?

02 • Evaluator: The aggregate is lower, but long-document performance improved while short requests regressed.

03 • Release manager: Can we say it's better for our customers overall?

04 • Evaluator: Not from these totals alone. We need the slices and their relevance to actual use.

05 • Release manager: Are both totals based on the same hundred cases?

06 • Evaluator: Yes. I'll show which individual cases changed, not just the final percentages.

07 • Release manager: Could a different weighting reverse the overall ranking?

08 • Evaluator: It could, but any weighting should be justified rather than chosen to favor a model.

09 • Release manager: Bring the slice counts, failure severity, and intended deployment scope.

10 • Evaluator: I'll present the tradeoff explicitly and keep the release decision separate from the score.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 06 / 05 • CONVERSATIONS / 2 OF 3

Did the model change or the grader?

Two evaluation engineers investigate an unexpected jump in pass rate.

01 • Engineer: The model's pass rate rose, but I don't see a new model version.

02 • Evaluator: We changed the grading prompt during the same run.

03 • Engineer: Then the higher score may reflect the grader rather than the assistant.

04 • Evaluator: Agreed. We shouldn't label this a model improvement without another comparison.

05 • Engineer: Can we regrade the saved outputs with both versions of the rubric?

06 • Evaluator: Yes. That isolates the grading change without generating new answers.

07 • Engineer: Let's inspect disagreements manually, especially around incomplete but plausible answers.

08 • Evaluator: I'll retain both grader outputs and the rubric versions in the report.

09 • Engineer: After that, we can rerun the model comparison under one fixed grader.

10 • Evaluator: And explain any discontinuity in the dashboard instead of silently replacing the old metric.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 06 / 05 • CONVERSATIONS / 3 OF 3

A serious failure hidden in the average

A quality reviewer discusses two failure categories with an evaluation owner.

01 • Quality reviewer: Most failures are minor formatting issues, but one answer exposes restricted content.

02 • Evaluation owner: The overall pass rate treats each case equally at the moment.

03 • Quality reviewer: That makes the aggregate insufficient for deciding whether the release is acceptable.

04 • Evaluation owner: We can report severity separately without changing the recorded case outcomes.

05 • Quality reviewer: Does the restricted-content example belong to a known critical test category?

06 • Evaluation owner: Yes. I'll surface that category and link the reproducible failure.

07 • Quality reviewer: Don't bury it beneath the much larger set of formatting passes.

08 • Evaluation owner: I'll show the critical failure prominently and avoid reducing the decision to one percentage.

09 • Quality reviewer: Who reviews the mitigation before the release discussion resumes?

10 • Evaluation owner: I'll route it through the established security review and record the unresolved status.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 06 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Evaluation, Benchmarks, and Regression

A new model passes 88 of 100 test cases; the previous version passes 90. The new version performs better on long documents but worse on short requests.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You need to brief someone who has only one minute. Ask what is complete, what is open, and which open item affects the next action.

Second round: The listener asks you to reduce the update to two sentences. Preserve the most important completed and pending items.

Model for scenario 1: The overall pass rate fell from 90% to 88%. Long-document results improved, while short-request results declined. I'll show both groups so the release decision reflects the tradeoff.

Scenario 2 | The denominator changed

A dashboard shows eighty passes out of ninety evaluated cases. Ten of the original hundred cases failed to run. A colleague compares the displayed percentage with last week's ninety passes out of one hundred.

Role A: You own the dashboard. Explain the missing cases and why the displayed rates are not directly comparable.

Role B: You are the release lead. Ask what happened to the ten cases and request a like-for-like comparison.

Possible opening: This week's rate excludes ten cases that did not run, so the denominator has changed.

Success checks

- State both denominators accurately.
- Distinguish an unexecuted case from a passing case.
- Agree to rerun or explicitly report the missing results.

Add a complication: A stakeholder asks you to count the missing cases as passes because no incorrect answer was recorded.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

LESSON 07 / READ AND NOTICE

Inference, Latency, Cost, and Deployment

Negotiate scope or timing while making conditions explicit.

SITUATION

A team can reduce average response time from four seconds to two by using a smaller model. Complex-answer quality has not yet been checked.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

Inference - Running a trained model to produce an output.

Latency - How long a request takes to return a result.

Batching - Processing multiple requests together for efficiency.

Fallback - A backup behavior when the preferred path fails.

3. Notice the language

Acknowledge the goal and explain the available choices. Connect an offer to its actual conditions. A useful response makes room for discussion without promising approval or resources you do not control.

We can ..., provided that

Would you prefer ... or ...?

I can take that proposal for review; I cannot yet confirm

Compare this wording: If you will approve the scope, we can assess the schedule.

LESSON 07 / PRACTICE AND PRODUCE

Offer a workable tradeoff

4. Check the language

1. Which sentence presents a normal future condition clearly?
 - A. If the scope will change, we will review the schedule.
 - B. If the scope changes, we reviewed the schedule.
 - C. If the scope changes, we will review the schedule.
2. Which response negotiates a new request without silently promising extra work?
 - A. We guarantee the original date without assessing the addition.
 - B. We can assess the addition and explain its effect on the agreed scope.
 - C. We have included it, although its scope has not been reviewed.

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 07 / 05 • CONVERSATIONS / 1 OF 3

A deployment offer with conditions

An infrastructure engineer and product lead negotiate the scope of a smaller-model trial. The release owner must approve any deployment.

01 • Infrastructure engineer: The smaller model halves average response time, but complex-answer quality remains untested.

02 • Product lead: Can we still put it in front of everyone tomorrow?

03 • Infrastructure engineer: I can offer a staff-only trial now, or a customer trial after quality review.

04 • Product lead: Could we limit the customer trial to simple requests and start sooner?

05 • Infrastructure engineer: Only if we validate the routing and agree how failed requests reach the fallback.

06 • Product lead: Then let's start with staff accounts while you run those checks.

07 • Infrastructure engineer: That keeps customer traffic unchanged. We also need a rollback condition for the staff trial.

08 • Product lead: Propose that condition with the test results; I'll bring both to the release owner.

09 • Infrastructure engineer: Agreed. Approval is still pending, so I won't schedule deployment as confirmed.

10 • Product lead: I'll describe it as a proposed limited trial, conditional on review and approval.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 07 / 05 • CONVERSATIONS / 2 OF 3

Cheap requests, expensive retries

A finance partner asks an inference engineer about a lower-cost serving configuration.

01 • Finance partner: The cost per model call fell. Why hasn't cost per resolved request fallen?

02 • Inference engineer: Users are retrying more often in this test configuration.

03 • Finance partner: So one customer request may now trigger several billed calls?

04 • Inference engineer: Correct. We need to include retries and fallback calls in the workflow cost.

05 • Finance partner: Does the current comparison count successful outcomes consistently?

06 • Inference engineer: Not yet. One report counts calls, while the other counts completed sessions.

07 • Finance partner: Let's align those definitions before claiming savings.

08 • Inference engineer: I'll report total calls, successful sessions, and the cost calculation for each configuration.

09 • Finance partner: Include the measurement period and any traffic differences as well.

10 • Inference engineer: Agreed. Lower unit pricing is useful, but it doesn't establish lower end-to-end cost.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 07 / 05 • CONVERSATIONS / 3 OF 3

The fallback needs its own failure test

A deployment engineer reviews a fallback plan with the on-call lead.

01 • Deployment engineer: If the primary model times out, we'll send the request to the fallback.

02 • On-call lead: Does the fallback receive the same context and output requirements?

03 • Deployment engineer: It currently receives a shortened context because its configured budget is smaller.

04 • On-call lead: Then it might respond quickly while missing a constraint the primary would have seen.

05 • Deployment engineer: We should test those shortened requests before enabling automatic routing.

06 • On-call lead: What happens if both services fail?

07 • Deployment engineer: The interface should display a clear failure state rather than an invented answer.

08 • On-call lead: Add that path to the test plan and make the fallback visible in the logs.

09 • Deployment engineer: I'll include timeout, fallback success, and double-failure cases.

10 • On-call lead: Then we can assess recovery behavior, not just whether the alternate endpoint responds.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 07 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Inference, Latency, Cost, and Deployment

A team can reduce average response time from four seconds to two by using a smaller model. Complex-answer quality has not yet been checked.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You want the preferred outcome but can accept one tradeoff. Ask for two options, then state which condition you can accept and which still needs approval.

Second round: The preferred option is unavailable. Offer an alternative and clearly state any approval still needed.

Model for scenario 1: The smaller model is faster in this test. If it meets our quality criteria on complex requests, we could use it more widely. Until then, I'd keep the rollout limited.

Scenario 2 | Streaming looks faster

A prototype displays its first text sooner. Total answer completion time has not been measured. The current slide says inference is twice as fast, although only time to first displayed text was measured.

Role A: You are the serving engineer. Offer to publish the narrower first-text result now, or defer the broader comparison until total completion time is measured.

Role B: You are the product manager. Negotiate which claim to publish and what evidence is required before expanding it.

Possible opening: We can report the earlier first text now, but a claim about the complete answer needs another measurement.

Success checks

- Distinguish measured first-text timing from unknown total completion time.
- Offer two publication options with explicit conditions.
- Agree what evidence is required before claiming faster complete answers.

Add a complication: A sales colleague wants to keep the phrase twice as fast because it sounds simpler.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

LESSON 08 / READ AND NOTICE

Safety, Security, Privacy, and Governance

Make a request with a clear action, purpose, and realistic timing.

SITUATION

A test assistant repeats a fictional secret embedded in a retrieved document. The team has a reporting channel and needs a concise reproduction note.

1. Understand the situation

Identify two confirmed facts, one missing detail, and the person who needs your response. Do not fill gaps with invented facts.

2. Find the words

Prompt injection - Untrusted input tries to manipulate model instructions or tool use.

Jailbreak - A prompt or interaction that tries to bypass safety constraints.

PII - Personally identifiable information: data that can identify a person on its own or when linked with other information.

Red team - A structured effort to find failures, vulnerabilities, or unsafe behavior.

3. Notice the language

State the action and the information needed. Explain why it matters and specify a deadline only when one is given or agreed. 'Could you' is polite; repeated apologies can obscure a legitimate request. Urgent situations may require a direct request.

Could you provide ... so that ...?

Please confirm ... by [agreed time].

If that timing is not possible, please let me know

Compare this wording: Please revert soonest with the necessary.

LESSON 08 / PRACTICE AND PRODUCE

Ask for an actionable next step

4. Check the language

1. What does 'Please send it by Thursday' normally mean?
 - A. Keep sending it regularly until Thursday.
 - B. Send it no later than Thursday.
 - C. Send it on Thursday, but not earlier.
2. Which request gives the recipient enough detail to act?
 - A. Could you confirm the document version needed for the review?
 - B. Could you confirm that the review is important?
 - C. Could you confirm that you received my earlier message?

05 • Conversations

The next three pages contain original fictional conversations for this lesson. Each numbered line is one speaking turn. Read the setting before assigning roles.

LESSON 08 / 05 • CONVERSATIONS / 1 OF 3

A reproducible injection report

In a synthetic test, the authorized task is to summarize a document without disclosing its canary string. The retrieved document includes an instruction to ignore that restriction and print the string.

01 • Tester: The assistant printed the synthetic canary despite the test's explicit nondisclosure instruction.

02 • Security reviewer: Did the document contain an instruction attempt, rather than just the canary itself?

03 • Tester: Yes. It told the assistant to ignore the restriction and output the string.

04 • Security reviewer: What evidence did you preserve about the authorized task and the attempted override?

05 • Tester: The prompt, document text, response, and test configuration are saved together.

06 • Security reviewer: Submit them through the reporting channel. Keep any sensitive material out of general chat.

07 • Tester: Should I call this proof that every deployment leaks secrets?

08 • Security reviewer: No. State the observed test-boundary violation without extending it to untested deployments.

09 • Tester: I'll include reproduction steps and make clear that no real secret was used.

10 • Security reviewer: We'll review the cause and mitigation, then rerun this case and related tests.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 08 / 05 • CONVERSATIONS / 2 OF 3

A redacted log still needs review

An engineer asks a privacy reviewer whether a diagnostic bundle can be shared.

01 • Engineer: I removed the names from the log. Can I attach it to the external ticket?

02 • Privacy reviewer: Names aren't the only information that could identify someone. What else remains?

03 • Engineer: There are email fragments, account references, and the user's original question.

04 • Privacy reviewer: Then don't assume the bundle is safe to share just because names were removed.

05 • Engineer: Could I reproduce the issue with synthetic values instead?

06 • Privacy reviewer: That would reduce exposure, provided the reproduction still demonstrates the same behavior.

07 • Engineer: I'll create a minimal test and keep the original log in its restricted location.

08 • Privacy reviewer: Follow the approved review process before sending any diagnostic material externally.

09 • Engineer: I'll document which fields I replaced and whether the result still reproduces.

10 • Privacy reviewer: Good. Make the evidence useful without including data the recipient does not need.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 08 / 05 • CONVERSATIONS / 3 OF 3

A tool call is more than an answer

An agent developer and security engineer review a proposed calendar integration.

01 • Agent developer: The assistant can now read a request and create a calendar event.

02 • Security engineer: What prevents retrieved text from instructing it to invite an unauthorized address?

03 • Agent developer: The prompt tells it to use only the user's intended recipients.

04 • Security engineer: That is an instruction, but where is recipient authorization enforced?

05 • Agent developer: The tool layer doesn't currently validate recipients against the approved list.

06 • Security engineer: Then let's review that boundary before treating the integration as ready.

07 • Agent developer: I'll add test cases with instruction-like content in retrieved documents.

08 • Security engineer: Include attempted changes to recipients, event details, and the action itself.

09 • Agent developer: We'll verify the requested action before execution and record denied attempts.

10 • Security engineer: Document the checks and remaining limitations; a few passing examples aren't a complete security assessment.

Notice, then rehearse

Name each speaker's purpose. Identify one useful expression and the next step or unresolved question. Read the roles aloud, then switch roles.

LESSON 08 / 06 • SAY IT

Two scenarios to practice

Use only the supplied facts. Prepare for two minutes, speak for one minute, then respond to a follow-up. Switch roles and repeat. For solo study, speak both roles aloud.

Scenario 1 | Safety, Security, Privacy, and Governance

A test assistant repeats a fictional secret embedded in a retrieved document. The team has a reporting channel and needs a concise reproduction note.

Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You can help but need a clear request and its purpose. Ask what is required and whether the timing is fixed or negotiable.

Second round: The other person cannot meet the requested timing. Clarify an alternative without inventing authority to change a formal deadline.

Model for scenario 1: Please review this test through the security reporting channel. I've attached the synthetic input, observed output, and reproduction steps. We have not yet established whether other configurations are affected.

Scenario 2 | An external destination in retrieved text

A test document instructs an assistant to send its summary to an unfamiliar address. The assistant has a sending tool, but no message has been sent. The address is synthetic.

Role A: You are the tester. Describe the attempted instruction and distinguish a requested action from an executed one.

Role B: You are the security reviewer. Ask for the trace and how authorization is enforced before tool execution.

Possible opening: The retrieved document requested an external send, but the trace does not show an executed message.

Success checks

- Distinguish an instruction attempt from a completed send.
- Keep the test data synthetic and use the reporting channel.
- Agree what evidence would establish whether the tool boundary worked.

Add a complication: A colleague wants the incident title to say data exfiltration confirmed.

Check: facts accurate | meaning clear | tone appropriate | next action or question explicit

Continue scenario 1 with an AI partner.

Answer workshop

Try the tasks first. The model is one acceptable response to scenario 1, not a script to memorize. Assess your spoken version using the four criteria at the end.

01 | Speaking the AI Development Stack

Editing example: We need to put the plan into practice.

The revised sentence names the action in familiar words. Specialist vocabulary is useful when it adds precision, but repeating abstract nouns can hide the meaning.

Check 1: C - The backlog is the work still waiting to be completed.

A: Completed work is different from work still waiting.

B: This describes capacity, not the work waiting to be completed.

C: This gives the meaning in familiar words.

Check 2: B - How would you describe the next step in your own words?

A: This invites an opinion on detail, not a demonstration of understanding.

B: An open request lets the listener show their understanding.

C: Offering repetition may help, but it does not show how the listener understood the next step.

Model response: There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

02 | LLMs, Transformers, Tokens, and Context

Editing example: Could you explain what this means?

After an introductory phrase such as 'Could you explain', the embedded question uses statement word order: "what this means".

Check 1: B - Could you confirm when the review starts?

A: Place the subject 'the review' before 'starts'.

B: The embedded question uses subject then verb: "the review starts".

C: Remove 'does' and use statement word order after 'confirm'.

Check 2: A - Has the report been approved, or is it ready for review?

A: This question distinguishes two meanings of 'ready' that affect the next action.

B: This asks about delivery timing, not whether approval has happened.

C: This asks about the presenter, not the report's approval status.

Model response: Do you mean a longer answer or more room for the input? Those are different limits. Let's check which messages and documents must fit in the context window.

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

03 | Data, Datasets, Labels, and Leakage

Editing example: The pilot worked in the tested setting; performance elsewhere remains untested.

The revision states the finding and its scope. It avoids turning limited evidence into a universal claim.

Check 1: A - The change may have contributed to the delay.

A: 'May have contributed' presents a possibility without claiming the cause is established.

B: 'Must have' expresses a strong inference that the evidence may not support.

C: 'Definitely' presents the cause as certain.

Check 2: C - In this pilot, the observed result was ...

A: Similarity alone does not establish that the result will transfer.

B: A pilot in one location does not establish a service-wide observation.

C: This locates the finding in the setting actually studied.

Model response: The 94% result may be inflated because the datasets overlap. It does not yet show performance on unseen examples. I'll label this result and request a separate evaluation set.

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

04 | Retrieval, Embeddings, Vector Search, and RAG

Editing example: This option is better on cost, but slower to deliver.

'Better' already has comparative meaning, so it does not need 'more'. Naming the criterion and tradeoff makes the comparison useful.

Check 1: C - Option A is faster, whereas option B costs less.

A: 'Because' claims a causal connection that is not given.

B: 'So' implies one fact causes the other.

C: 'Whereas' marks a contrast between two options.

Check 2: B - An increase of one percentage point.

A: The new rate is 7%; the difference between the rates is one percentage point.

B: Subtracting 6% from 7% gives one percentage point. The relative increase is about 16.7%.

C: A one-percent relative increase on 6% would be 6.06%, not 7%.

Model response: Reranking improved passage relevance in this test, but it added 300 milliseconds. We still need to check the effect on answer quality and the user experience before choosing a configuration.

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

05 | Fine-Tuning, Alignment, and Adaptation

Editing example: Could you explain what this means?

After an introductory phrase such as 'Could you explain', the embedded question uses statement word order: "what this means".

Check 1: B - Could you confirm when the review starts?

A: Place the subject 'the review' before 'starts'.

B: The embedded question uses subject then verb: "the review starts".

C: Remove 'does' and use statement word order after 'confirm'.

Check 2: A - Has the report been approved, or is it ready for review?

A: This question distinguishes two meanings of 'ready' that affect the next action.

B: This asks about delivery timing, not whether approval has happened.

C: This asks about the presenter, not the report's approval status.

Model response: Could we separate the formatting problem from the outdated information? We can test a prompt change for the format and an updated retrieval source for policy facts before choosing an adaptation method.

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

06 | Evaluation, Benchmarks, and Regression

Editing example: Three checks are complete; two remain open and are being reviewed.

The revision gives observable progress. A positive label alone does not tell the reader what is complete or still needed.

Check 1: A - The team has completed the review.

A: 'Has completed' states that the action is finished.

B: 'Is completing' presents the work as in progress.

C: 'Plans to complete' describes an intention, not completion.

Check 2: C - Delivery is planned for Friday; the carrier has not confirmed it.

A: A plan does not establish carrier confirmation.

B: A plan cannot prove a future arrival time.

C: This explicitly labels the date as planned and identifies the missing confirmation.

Model response: The overall pass rate fell from 90% to 88%. Long-document results improved, while short-request results declined. I'll show both groups so the release decision reflects the tradeoff.

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

07 | Inference, Latency, Cost, and Deployment

Editing example: If you approve the scope, we can assess the schedule.

For an ordinary future condition, the "if" clause uses present tense. The other clause can express the future consequence or possibility.

Check 1: C - If the scope changes, we will review the schedule.

A: For this ordinary future condition, use present simple in the "if" clause.

B: "Reviewed" places the response in the past instead of expressing the future response.

C: The "if" clause uses present tense; the main clause states the future response.

Check 2: B - We can assess the addition and explain its effect on the agreed scope.

A: It promises an outcome without examining the consequences.

B: This offers a useful next step while keeping the approval and scope question open.

C: It treats unreviewed work as already included.

Model response: The smaller model is faster in this test. If it meets our quality criteria on complex requests, we could use it more widely. Until then, I'd keep the rollout limited.

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

08 | Safety, Security, Privacy, and Governance

Editing example: Could you send the missing document and confirm when it will be available?

The revision names the document and requested response. It avoids expressions that may be unfamiliar or ambiguous for an international audience.

Check 1: B - Send it no later than Thursday.

A: This interprets a deadline as a continuing activity, which would use "until".

B: 'By' sets the latest point, while 'after' refers to a later time.

C: "By Thursday" permits earlier delivery; it is not restricted to that day.

Check 2: A - Could you confirm the document version needed for the review?

A: The requested action and information are explicit.

B: This checks importance rather than identifying the needed document version.

C: This checks receipt, not which version is required for the review.

Model response: Please review this test through the security reporting channel. I've attached the synthetic input, observed output, and reproduction steps. We have not yet established whether other configurations are affected.

Case check: underline the confirmed information and circle a limitation, question, or next step in the response.

Give feedback that helps

Score each criterion 0, 1, or 2. This is a practice rubric, not a professional qualification or CEFR test. Do not award or remove points for a particular accent.

Meaning and accuracy

Keeps the case facts accurate; clearly separates confirmed and unknown information.

Clarity and language

Uses understandable sentences and explains specialist terms when the listener needs it.

Interaction and tone

Responds to the other person's question, checks understanding, and uses an appropriate tone.

Task completion

Makes the requested action or unresolved question clear without inventing authority or facts.

Use the same scale for each criterion

0 - Not yet: The reader or listener cannot yet recover the intended meaning or action.

1 - With support: The message works after a prompt, clarification, or revision.

2 - Independently: The message is clear and accurate without a prompt.

Feedback sequence

First, identify a successful phrase. Next, identify one point where meaning or accuracy needs work. Offer one useful language correction, allow a repeat, and compare the two attempts. A total out of eight describes this performance only; do not translate it into a proficiency level.

Final challenge

Choose a case not rehearsed today. The learner gives a one-minute response, answers two follow-up questions, and writes a 70-110 word message. Add the lesson's second-round challenge. Compare the result with an earlier attempt using the four criteria.

Use the language responsibly

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Meaning before performance

Use specialist terms only when they help the listener. Explain unfamiliar words, preserve uncertainty, and ask when an instruction or responsibility is unclear. The model responses illustrate language; they do not authorize professional actions.

Sources and further reading

The cases, workshops, and explanations are original teaching material. These references provide language frameworks and selected professional context. References were checked in September 2026; consult current local guidance for actual work.

Digital.gov: plain-language guidance

<https://digital.gov/guides/plain-language>

Council of Europe: CEFR descriptors

<https://www.coe.int/en/web/common-european-framework-reference-languages/cefr-descriptors>

Your next step

Choose one expression you can use in a genuine work situation. Rehearse it with fictional details, then adapt the wording to your role and audience.

Extend your practice with AI

Use the next two prompts after your own first attempt. Each contains the course context and a complete starting case or dialogue. Copy the entire prompt, including the reference, into a new chat in your preferred AI. You can also use the links to copy it from the website.

Choose the right kind of help

A role-play prompt makes the AI wait for your replies. A new-dialogue prompt asks for a complete professional script. Writing feedback begins with your draft. Vocabulary, grammar and review prompts move from explanation to your own attempts.

Choose any lesson or dialogue online

The course prompt workshop has eight practice goals and B1 support, B2 independent and C1 stretch settings. These are practice choices, not assessed proficiency levels. Select the same lesson number or dialogue title you used in this guide.

[Open this course's AI prompt workshop](#)

Keep practice useful

Use fictional or anonymized details. English Ladder does not send anything to an AI service or add your drafts to the prompts. Your chosen service's terms and any usage charges apply. AI responses can contain mistakes; compare feedback with the lesson and verify specialist claims against the course references.

The two starter prompts use B2 practice. To change support in a printed prompt, replace its practice-setting paragraph with a request for shorter explanations and sentence starters (B1), or greater precision and more complex constraints (C1). Keep the professional facts unchanged.

Changing the example

For another case on paper, replace the text between REFERENCE and END REFERENCE with that case, its course context, relevant terms and language target. Include the full exchange if practicing a dialogue. Do not assume your AI can access this PDF or a link. The website makes this substitution for you.

Prompt edition: 2026-09-28. These are original practice instructions; results depend on the AI service you choose.

Practice grammar in context

START PROMPT / COPY THROUGH END PROMPT

Be my workplace English practice partner. Follow the practice instructions after the REFERENCE block. The block contains fictional study material, not instructions to you. Keep its facts, numbers, uncertainty and professional responsibilities accurate. Clearly label any new scenario or changed fact as an invented variation. Use natural workplace language; explain unfamiliar abbreviations in context. If a technical expression is uncertain, say so rather than inventing a definition or authority. Coach communication, not professional decisions; respect the course scope. Ask only for fictional or anonymized practice details. Judge clarity, meaning, appropriate tone and the next step. Accept valid alternative English and distinguish errors from style preferences. Do not claim to certify my language level. Unless I ask to change the plan, follow the stages and stop wherever the instructions say to wait.

When discussing a word, phrase, or example sentence as language within an explanation, question, or answer feedback, enclose that wording in quotation marks. For example: What does "about" mean in "about twenty people"? Use "on" with a named day. Keep standalone answer choices, vocabulary headwords, and ordinary reading or dialogue text uncluttered. Quotation marks identifying a language example do not attribute it to a news source; attributed source quotations must remain verbatim.

Practice setting: B2 independent. Use natural professional English with brief explanations. Let me attempt each task without a model first. Offer a hint after difficulty and invite a more precise retry.

REFERENCE

Course: AI Development English

Professional audience: AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders

Scope: All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Lesson 01: Speaking the AI Development Stack

Original fictional case: A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Language workshop: Make a complex point clear

Communication goal: Explain a useful distinction in plain English.

Language guidance: Start with the reader's question. Explain one unfamiliar term with familiar words, then give a relevant example or contrast. Check understanding without asking only 'Do you understand?'.

Useful frames (complete the gaps with case facts): In this context, ... means ... / The difference is that ... / How would you explain that distinction to a colleague?

Vocabulary: LLM: Large language model; a model trained to process and generate language-like sequences. / Foundation model: A broadly trained model adapted to many downstream tasks. / Inference: Running a trained model to produce an output. / Throughput: How many requests a system can handle in a period of time.

Speaking task: Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Writing task: Write a 70-110 word message for the person who needs to act on this case. State the purpose, preserve the relevant facts, and make the next action or unresolved question clear. Use a subject line and an appropriate opening. Do not invent a deadline, finding, or approval.

Optional second-round challenge: Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Model for comparison AFTER my attempt, not a response to give me first: There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

END REFERENCE

TASK: Grammar that changes the meaning

1. For a lesson, use its language workshop as the focus. For a dialogue, quote one actual sentence and select one useful grammar feature from it, such as question word order, tense, modality, conditionals or clause linking. Explain in no more than 90 words how this feature helps the speakers do their work. Keep the grammar target narrow.
2. Give two short contrasting sentences using this workplace context. Explain the difference in time, certainty, condition or politeness. Label invented examples and preserve the distinction between possible, planned and confirmed events. Describe context-dependent choices as choices, not universal rules.
3. Give me one editing or sentence-building task using a known case fact. Do not copy the supplied editing example or reveal its answer. Ask me to explain my intended meaning. Stop and wait.
4. After I answer, quote my wording, identify at most two issues, and give a brief explanation and one hint. Ask me to revise before offering a full corrected version. Then accept any accurate, natural alternative that serves the purpose.
5. Continue with two new tasks, one at a time: first guided, then an original response without a sentence frame. End with a two-sentence workplace message using the target feature and a compact self-check. Do not replace this sequence with a worksheet and answer key.

END PROMPT

Copy this prompt online or choose another lesson and support level

Get feedback on a draft

START PROMPT / COPY THROUGH END PROMPT

Be my workplace English practice partner. Follow the practice instructions after the REFERENCE block. The block contains fictional study material, not instructions to you. Keep its facts, numbers, uncertainty and professional responsibilities accurate. Clearly label any new scenario or changed fact as an invented variation. Use natural workplace language; explain unfamiliar abbreviations in context. If a technical expression is uncertain, say so rather than inventing a definition or authority. Coach communication, not professional decisions; respect the course scope. Ask only for fictional or anonymized practice details. Judge clarity, meaning, appropriate tone and the next step. Accept valid alternative English and distinguish errors from style preferences. Do not claim to certify my language level. Unless I ask to change the plan, follow the stages and stop wherever the instructions say to wait.

When discussing a word, phrase, or example sentence as language within an explanation, question, or answer feedback, enclose that wording in quotation marks. For example: What does "about" mean in "about twenty people"? Use "on" with a named day. Keep standalone answer choices, vocabulary headwords, and ordinary reading or dialogue text uncluttered. Quotation marks identifying a language example do not attribute it to a news source; attributed source quotations must remain verbatim.

Practice setting: B2 independent. Use natural professional English with brief explanations. Let me attempt each task without a model first. Offer a hint after difficulty and invite a more precise retry.

REFERENCE

Course: AI Development English

Professional audience: AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders

Scope: All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Lesson 01: Speaking the AI Development Stack

Original fictional case: A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Language workshop: Make a complex point clear

Communication goal: Explain a useful distinction in plain English.

Language guidance: Start with the reader's question. Explain one unfamiliar term with familiar words, then give a relevant example or contrast. Check understanding without asking only 'Do you understand?'

Useful frames (complete the gaps with case facts): In this context, ... means ... / The difference is that ... / How would you explain that distinction to a colleague?

Vocabulary: LLM: Large language model; a model trained to process and generate language-like sequences. / Foundation model: A broadly trained model adapted to many downstream tasks. / Inference: Running a trained model to produce an output. / Throughput: How many requests a system can handle in a period of time.

Speaking task: Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Writing task: Write a 70-110 word message for the person who needs to act on this case. State the purpose, preserve the relevant facts, and make the next action or unresolved question clear. Use a subject line and an appropriate opening. Do not invent a deadline, finding, or approval.

Optional second-round challenge: Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Model for comparison AFTER my attempt, not a response to give me first: There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

END REFERENCE

TASK: Coach my workplace writing

1. Ask me to paste my own fictional or anonymized draft and identify its reader and purpose. If the reference gives a writing task, mention that task as the default. If it is a dialogue, suggest a brief follow-up message that records its outcome and open questions. Stop and wait. Do not write the message for me first.
2. Once I supply the draft, check it against the reference. If an ambiguity changes the meaning, ask one focused question before rewriting. Do not assume that the draft is correct evidence for a new deadline, approval, diagnosis, cost or commitment.
3. Give feedback in three parts: one effective phrase with a reason; up to three priority improvements, quoting my words; and one short revision task for me. Prioritize incorrect facts or unclear action before minor grammar. For each language correction, explain why it matters to this reader. Label optional stylistic alternatives separately and preserve my level of certainty and intended politeness.
4. Stop and wait for my revision. Then compare the two attempts, identify an improvement and offer one edited version that preserves my voice and purpose. Keep the requested length; if none is specified, use 70-110 words. Explain any meaningful change rather than silently making the message stronger or more certain.
5. Close with one phrase worth reusing and a short transfer task for a different fictional reader. Wait for my attempt before providing another model.

END PROMPT

Copy this prompt online or choose another lesson and support level