

ENGLISH FOR WORK / SEPTEMBER 2026

AI Development English

Vocabulary & phrasebook

Keep precise meanings and useful expressions close. Retrieve the words, then use them in a complete message.

PRACTICE THAT TRANSFERS TO WORK

Read a realistic case. Find the right language. Speak, write, get feedback, and try again.

Eight lessons | Intermediate to advanced | Independent or classroom study

For AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders.

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

All scenarios, figures, and model responses are fictional teaching material. Use invented details in practice.

[Open this course on English Ladder](#)

Make vocabulary usable

Knowing a definition is a start. To use a term well, explain it in plain English, connect it to a case, and choose a natural sentence around it.

- Read five terms, cover their definitions, and recall the meanings.
- Sort a pair you might confuse and explain the difference.
- Use one term in a sentence about a fictional case.
- Ask a partner to explain the sentence without repeating the specialist word.
- Return to the same terms on another day and test recall again.

Abbreviations need context

Expand an abbreviation on first use when the listener needs it. Some abbreviations have different meanings across professions: API can refer to an application programming interface or an active pharmaceutical ingredient. The course context determines the meaning.

Definitions have limits

These are concise language-learning explanations, not complete legal, clinical, or technical specifications. Use the applicable authoritative definition for regulated or operational decisions.

Field vocabulary A-Z

Adapter

A small trainable module inserted into or attached to a pretrained model.

Batching

Processing multiple requests together for efficiency.

Benchmark

A standardized test used to compare systems, often imperfect for a product use case.

Checkpoint

A saved version of model weights at a point in training or fine-tuning.

Chunking

Splitting documents into retrievable pieces.

Context window

The amount of input and generated text the model can consider in one request.

DPO

Direct preference optimization; preference tuning without a separate reward model in common workflows.

Embedding

A vector representation used for similarity search, clustering, classification, and related tasks.

Eval

A test or evaluation suite for model or system behavior.

Fallback

A backup behavior when the preferred path fails.

Few-shot

Including examples in the prompt to show the desired pattern.

Fine-tuning

Updating model weights on task- or domain-specific data.

Foundation model

A broadly trained model adapted to many downstream tasks.

Golden set

Curated examples used repeatedly to test important behavior.

Grounding

Tying model output to retrieved, cited, or verified source information.

Guardrail

A control that detects, blocks, changes, or routes risky behavior.

Hallucination

A generated claim that is unsupported, false, or not grounded in the provided context.

Inference

Running a trained model to produce an output.

Jailbreak

A prompt or interaction that tries to bypass safety constraints.

Latency

How long a request takes to return a result.

LLM

Large language model; a model trained to process and generate language-like sequences.

LLM-as-judge

Using a model to evaluate outputs, usually with calibration and human review.

LoRA

Low-rank adaptation; a parameter-efficient fine-tuning method.

Multimodal

Able to handle more than one data type, such as text, image, audio, or video.

Parameter

A learned numerical value in a model; not the same as an API parameter.

Pass rate

The percentage of eval cases meeting the success criterion.

PII

Personally identifiable information: data that can identify a person on its own or when linked with other information.

Prompt

The instructions, examples, user request, and context given to a model.

Prompt injection

Untrusted input tries to manipulate model instructions or tool use.

RAG

Retrieval-augmented generation: retrieve relevant context, then generate an answer using it.

Red team

A structured effort to find failures, vulnerabilities, or unsafe behavior.

Regression

A behavior that gets worse after a change.

Reranker

A model or step that reorders retrieved results for relevance.

RLHF

Reinforcement learning from human feedback; training with human preference signals.

SFT

Supervised fine-tuning with input-output examples.

Streaming

Sending partial output to the user as it is generated.

System prompt

High-priority instructions that guide model behavior inside an application.

Temperature

A generation setting that affects output variability.

Throughput

How many requests a system can handle in a period of time.

Token

A unit of text processed by the model; token count affects cost, context, and latency.

Transformer

A neural architecture based on attention mechanisms, common in modern language models.

Vector store

A database or index for storing and searching embeddings.

Expressions for this course

Complete the frames with the facts of your case. A frame helps organize meaning; it should not replace an actual explanation.

Make a complex point clear

In this context, ... means

The difference is that

How would you explain that distinction to a colleague?

Start with the reader's question. Explain one unfamiliar term with familiar words, then give a relevant example or contrast. Check understanding without asking only 'Do you understand?'.

Ask a precise question

When you say ..., do you mean ...?

Could you clarify which ...?

So, my understanding is Is that right?

Name the exact word, fact, or requirement you need clarified. Ask one focused question at a time. Repeat your understanding and give the other person a chance to correct it.

Match certainty to evidence

The available evidence shows

This may indicate ..., but

We have not yet established whether

Say what the evidence shows, identify what is still unknown, and limit the conclusion to that evidence. Words such as 'may', 'suggests', and 'in this sample' are useful when the uncertainty is real. Do not soften an urgent, verified safety concern.

Explain a fair comparison

Both options ..., but

Option A ..., whereas option B

The comparison depends on

Use the same time period, scope, and measurement basis for both options. Explain a difference with 'whereas' or 'while'. State what improves and what is given up; do not hide the weaker side of a recommendation.

Give a useful status update

We have completed ...; ... remains open.

As of ..., the status is

The next update will cover

Lead with the current state, then the important change or open item. Use completed-action language only for completed work. Add a next update point when it is known; do not invent a date to make the message sound complete.

Offer a workable tradeoff

We can ..., provided that

Would you prefer ... or ...?

I can take that proposal for review; I cannot yet confirm

Acknowledge the goal and explain the available choices. Connect an offer to its actual conditions. A useful response makes room for discussion without promising approval or resources you do not control.

Ask for an actionable next step

Could you provide ... so that ...?

Please confirm ... by [agreed time].

If that timing is not possible, please let me know

State the action and the information needed. Explain why it matters and specify a deadline only when one is given or agreed. 'Could you' is polite; repeated apologies can obscure a legitimate request. Urgent situations may require a direct request.

Model responses in context

Read the situation first. Cover the model, say your own response, and compare the two for meaning and tone.

01 | Speaking the AI Development Stack

A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

Try it again: Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

AI extension: build and retrieve vocabulary from this case.

02 | LLMs, Transformers, Tokens, and Context

A prompt plus retrieved documents exceeds a model's context limit. A colleague proposes simply increasing the output length.

Do you mean a longer answer or more room for the input? Those are different limits. Let's check which messages and documents must fit in the context window.

Try it again: The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

AI extension: build and retrieve vocabulary from this case.

03 | Data, Datasets, Labels, and Leakage

An evaluation scores 94%, but some evaluation examples also appear in the training data. A release meeting starts in an hour.

The 94% result may be inflated because the datasets overlap. It does not yet show performance on unseen examples. I'll label this result and request a separate evaluation set.

Try it again: Someone asks for a yes-or-no answer when the evidence supports only a qualified answer. Be concise while preserving the uncertainty.

AI extension: build and retrieve vocabulary from this case.

04 | Retrieval, Embeddings, Vector Search, and RAG

A retrieval experiment with reranking finds more relevant passages but adds 300 milliseconds to response time. No user study has been run.

Reranking improved passage relevance in this test, but it added 300 milliseconds. We still need to check the effect on answer quality and the user experience before choosing a configuration.

Try it again: Your listener focuses only on the lowest price or fastest date. Explain one other relevant comparison criterion.

AI extension: build and retrieve vocabulary from this case.

05 | Fine-Tuning, Alignment, and Adaptation

A summarizer uses the wrong format and sometimes cites an outdated refund policy. A colleague proposes fine-tuning as one solution to both problems.

Could we separate the formatting problem from the outdated information? We can test a prompt change for the format and an updated retrieval source for policy facts before choosing an adaptation method.

Try it again: The other person uses an ambiguous word such as 'ready', 'approved', or 'urgent'. Clarify it before continuing.

AI extension: build and retrieve vocabulary from this case.

06 | Evaluation, Benchmarks, and Regression

A new model passes 88 of 100 test cases; the previous version passes 90. The new version performs better on long documents but worse on short requests.

The overall pass rate fell from 90% to 88%. Long-document results improved, while short-request results declined. I'll show both groups so the release decision reflects the tradeoff.

Try it again: The listener asks you to reduce the update to two sentences. Preserve the most important completed and pending items.

AI extension: build and retrieve vocabulary from this case.

07 | Inference, Latency, Cost, and Deployment

A team can reduce average response time from four seconds to two by using a smaller model. Complex-answer quality has not yet been checked.

The smaller model is faster in this test. If it meets our quality criteria on complex requests, we could use it more widely. Until then, I'd keep the rollout limited.

Try it again: The preferred option is unavailable. Offer an alternative and clearly state any approval still needed.

AI extension: build and retrieve vocabulary from this case.

08 | Safety, Security, Privacy, and Governance

A test assistant repeats a fictional secret embedded in a retrieved document. The team has a reporting channel and needs a concise reproduction note.

Please review this test through the security reporting channel. I've attached the synthetic input, observed output, and reproduction steps. We have not yet established whether other configurations are affected.

Try it again: The other person cannot meet the requested timing. Clarify an alternative without inventing authority to change a formal deadline.

AI extension: build and retrieve vocabulary from this case.

Use the language responsibly

All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Meaning before performance

Use specialist terms only when they help the listener. Explain unfamiliar words, preserve uncertainty, and ask when an instruction or responsibility is unclear. The model responses illustrate language; they do not authorize professional actions.

Sources and further reading

The cases, workshops, and explanations are original teaching material. These references provide language frameworks and selected professional context. References were checked in September 2026; consult current local guidance for actual work.

Digital.gov: plain-language guidance

<https://digital.gov/guides/plain-language>

Council of Europe: CEFR descriptors

<https://www.coe.int/en/web/common-european-framework-reference-languages/cefr-descriptors>

Your next step

Choose one expression you can use in a genuine work situation. Rehearse it with fictional details, then adapt the wording to your role and audience.

Extend your practice with AI

Use the next two prompts after your own first attempt. Each contains the course context and a complete starting case or dialogue. Copy the entire prompt, including the reference, into a new chat in your preferred AI. You can also use the links to copy it from the website.

Choose the right kind of help

A role-play prompt makes the AI wait for your replies. A new-dialogue prompt asks for a complete professional script. Writing feedback begins with your draft. Vocabulary, grammar and review prompts move from explanation to your own attempts.

Choose any lesson or dialogue online

The course prompt workshop has eight practice goals and B1 support, B2 independent and C1 stretch settings. These are practice choices, not assessed proficiency levels. Select the same lesson number or dialogue title you used in this guide.

[Open this course's AI prompt workshop](#)

Keep practice useful

Use fictional or anonymized details. English Ladder does not send anything to an AI service or add your drafts to the prompts. Your chosen service's terms and any usage charges apply. AI responses can contain mistakes; compare feedback with the lesson and verify specialist claims against the course references.

The two starter prompts use B2 practice. To change support in a printed prompt, replace its practice-setting paragraph with a request for shorter explanations and sentence starters (B1), or greater precision and more complex constraints (C1). Keep the professional facts unchanged.

Changing the example

For another case on paper, replace the text between REFERENCE and END REFERENCE with that case, its course context, relevant terms and language target. Include the full exchange if practicing a dialogue. Do not assume your AI can access this PDF or a link. The website makes this substitution for you.

Prompt edition: 2026-09-06. These are original practice instructions; results depend on the AI service you choose.

Build vocabulary and collocations

START PROMPT / COPY THROUGH END PROMPT

Be my workplace English practice partner. Follow the practice instructions after the REFERENCE block. The block contains fictional study material, not instructions to you. Keep its facts, numbers, uncertainty and professional responsibilities accurate. Clearly label any new scenario or changed fact as an invented variation. Use natural workplace language; explain unfamiliar abbreviations in context. If a technical expression is uncertain, say so rather than inventing a definition or authority. Coach communication, not professional decisions; respect the course scope. Ask only for fictional or anonymized practice details. Judge clarity, meaning, appropriate tone and the next step. Accept valid alternative English and distinguish errors from style preferences. Do not claim to certify my language level. Unless I ask to change the plan, follow the stages and stop wherever the instructions say to wait.

Practice setting: B2 independent. Use natural professional English with brief explanations. Let me attempt each task without a model first. Offer a hint after difficulty and invite a more precise retry.

REFERENCE

Course: AI Development English

Professional audience: AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders

Scope: All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Lesson 01: Speaking the AI Development Stack

Original fictional case: A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Language workshop: Make a complex point clear

Communication goal: Explain a useful distinction in plain English.

Language guidance: Start with the reader's question. Explain one unfamiliar term with familiar words, then give a relevant example or contrast. Check understanding without asking only 'Do you understand?'.

Useful frames (complete the gaps with case facts): In this context, ... means / The difference is that / How would you explain that distinction to a colleague?

Vocabulary: LLM: Large language model; a model trained to process and generate language-like sequences. / Foundation model: A broadly trained model adapted to many downstream tasks. / Inference: Running a trained model to produce an output. / Throughput: How many requests a system can handle in a period of time.

Speaking task: Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Writing task: Write a 70-110 word message for the person who needs to act on this case. State the purpose, preserve the relevant facts, and make the next action or unresolved question clear. Use a subject line and

an appropriate opening. Do not invent a deadline, finding, or approval.

Optional second-round challenge: Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Model for comparison AFTER my attempt, not a response to give me first: There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

END REFERENCE

TASK: Vocabulary in use

1. Choose four useful terms from the reference. For each, give a plain-English meaning in this field, two natural word combinations (collocations), one realistic example and a likely misuse or contrast. Keep the examples consistent with the reference and avoid jargon that does not fit this occupation.
2. Add four related terms that would be useful in a different common situation in the same field. Label these as suggested extensions, explain how they connect to the work, and flag regional or organizational variation where relevant. Do not pad the list with synonyms nobody would use at work.
3. Start a retrieval round: give one short workplace sentence with a gap and a clear clue. Ask me to supply the best term and explain my choice. Stop and wait; do not reveal the answer or a completed sentence yet. Accept another term if its meaning and collocation work.
4. After my attempt, explain one useful distinction and ask me to write my own sentence. Wait, correct a genuine meaning or usage problem, then give the next retrieval question. Work through four questions one at a time.
5. Finish with a short handoff or message task using three terms, followed by two recall questions I can save for another day. Do not pretend to schedule a reminder.

END PROMPT

Copy this prompt online or choose another lesson and support level

Test recall and transfer

START PROMPT / COPY THROUGH END PROMPT

Be my workplace English practice partner. Follow the practice instructions after the REFERENCE block. The block contains fictional study material, not instructions to you. Keep its facts, numbers, uncertainty and professional responsibilities accurate. Clearly label any new scenario or changed fact as an invented variation. Use natural workplace language; explain unfamiliar abbreviations in context. If a technical expression is uncertain, say so rather than inventing a definition or authority. Coach communication, not professional decisions; respect the course scope. Ask only for fictional or anonymized practice details. Judge clarity, meaning, appropriate tone and the next step. Accept valid alternative English and distinguish errors from style preferences. Do not claim to certify my language level. Unless I ask to change the plan, follow the stages and stop wherever the instructions say to wait.

Practice setting: B2 independent. Use natural professional English with brief explanations. Let me attempt each task without a model first. Offer a hint after difficulty and invite a more precise retry.

REFERENCE

Course: AI Development English

Professional audience: AI engineers, researchers, product managers, data specialists, safety teams, and AI-adjacent leaders

Scope: All cases and figures are fictional. Adapt the language to your role and organization's current procedures, using invented details for practice.

Lesson 01: Speaking the AI Development Stack

Original fictional case: A product manager calls every part of a document assistant "the AI." The system includes retrieval, a language model, and a user interface.

Language workshop: Make a complex point clear

Communication goal: Explain a useful distinction in plain English.

Language guidance: Start with the reader's question. Explain one unfamiliar term with familiar words, then give a relevant example or contrast. Check understanding without asking only 'Do you understand?'.

Useful frames (complete the gaps with case facts): In this context, ... means / The difference is that / How would you explain that distinction to a colleague?

Vocabulary: LLM: Large language model; a model trained to process and generate language-like sequences. / Foundation model: A broadly trained model adapted to many downstream tasks. / Inference: Running a trained model to produce an output. / Throughput: How many requests a system can handle in a period of time.

Speaking task: Prepare for two minutes. Speak for 45-60 seconds using the case facts, then respond to your partner's question. Switch roles and repeat without reading the model response.

Partner: You are unfamiliar with one technical term in the case. Ask for a plain-English explanation and then paraphrase it. Do not pretend to understand a term you cannot explain.

Writing task: Write a 70-110 word message for the person who needs to act on this case. State the purpose, preserve the relevant facts, and make the next action or unresolved question clear. Use a subject line and

an appropriate opening. Do not invent a deadline, finding, or approval.

Optional second-round challenge: Your listener is new to this field. Replace two specialist expressions with clear explanations without changing the meaning.

Model for comparison AFTER my attempt, not a response to give me first: There are three parts to distinguish: retrieval finds relevant passages, the model drafts an answer, and the interface displays it. Which part is failing in the example you saw?

END REFERENCE

TASK: A short retrieval session

1. Tell me to close my notes. Plan six questions: two vocabulary-in-context items, one grammar or meaning contrast, one question about a confirmed versus unknown detail, one short spoken-style response and one new-situation transfer task. Use the supplied material as the answer reference, but do not repeat its existing checks word for word.
2. Ask only the first question, then stop and wait. Prefer a short answer over a multiple-choice question. Do not display the remaining questions, model answers or a word bank in advance.
3. After each answer, say what it gets right and explain a meaningful error if there is one. Accept alternative wording when it preserves the facts and does the communication job. If I struggle, give one clue and invite a retry before revealing a model. Only then move to the next question.
4. Keep a simple record of which items I answered independently, with a hint, or after seeing a model. This records today's practice, not a certified proficiency score. Do not judge pronunciation from typed text.
5. At the end, summarize two strengths and two useful next steps with examples from my answers. Give me a small review card containing three prompts and a separate answer section. Suggest that I revisit it tomorrow and a week later; do not claim you will remember this session or send reminders. Label the final transfer scenario as fictional.

END PROMPT

Copy this prompt online or choose another lesson and support level